

*std::to\_string*  
faster than light

Афанасьев Иван  
cpp sberia 2020

# Intro

- *std::to\_string* / *std::to\_wstring*
- integers only
- libc++
- <https://reviews.llvm.org/D59178>

## std::to\_string

Defined in header `<string>`

<code>std::string to_string( int value );</code>	(1)	(since C++11)
<code>std::string to_string( long value );</code>	(2)	(since C++11)
<code>std::string to_string( long long value );</code>	(3)	(since C++11)
<code>std::string to_string( unsigned value );</code>	(4)	(since C++11)
<code>std::string to_string( unsigned long value );</code>	(5)	(since C++11)
<code>std::string to_string( unsigned long long value );</code>	(6)	(since C++11)
<code>std::string to_string( float value );</code>	(7)	(since C++11)
<code>std::string to_string( double value );</code>	(8)	(since C++11)
<code>std::string to_string( long double value );</code>	(9)	(since C++11)

Converts a numeric value to `std::string`.

# Intro

- *std::to\_string* / *std::to\_wstring*
- integers only
- libc++
- <https://reviews.llvm.org/D59178>

## std::to\_string

Defined in header `<string>`

<code>std::string to_string( int value );</code>	(1)	(since C++11)
<code>std::string to_string( long value );</code>	(2)	(since C++11)
<code>std::string to_string( long long value );</code>	(3)	(since C++11)
<code>std::string to_string( unsigned value );</code>	(4)	(since C++11)
<code>std::string to_string( unsigned long value );</code>	(5)	(since C++11)
<code>std::string to_string( unsigned long long value );</code>	(6)	(since C++11)
<code>std::string to_string( float value );</code>	(7)	(since C++11)
<code>std::string to_string( double value );</code>	(8)	(since C++11)
<code>std::string to_string( long double value );</code>	(9)	(since C++11)

Converts a numeric value to `std::string`.

- 1) Converts a signed integer to a string with the same content as what `std::sprintf(buf, "%d", value)` would produce for sufficiently large buf.
- 2) Converts a signed integer to a string with the same content as what `std::sprintf(buf, "%ld", value)` would produce for sufficiently large buf.

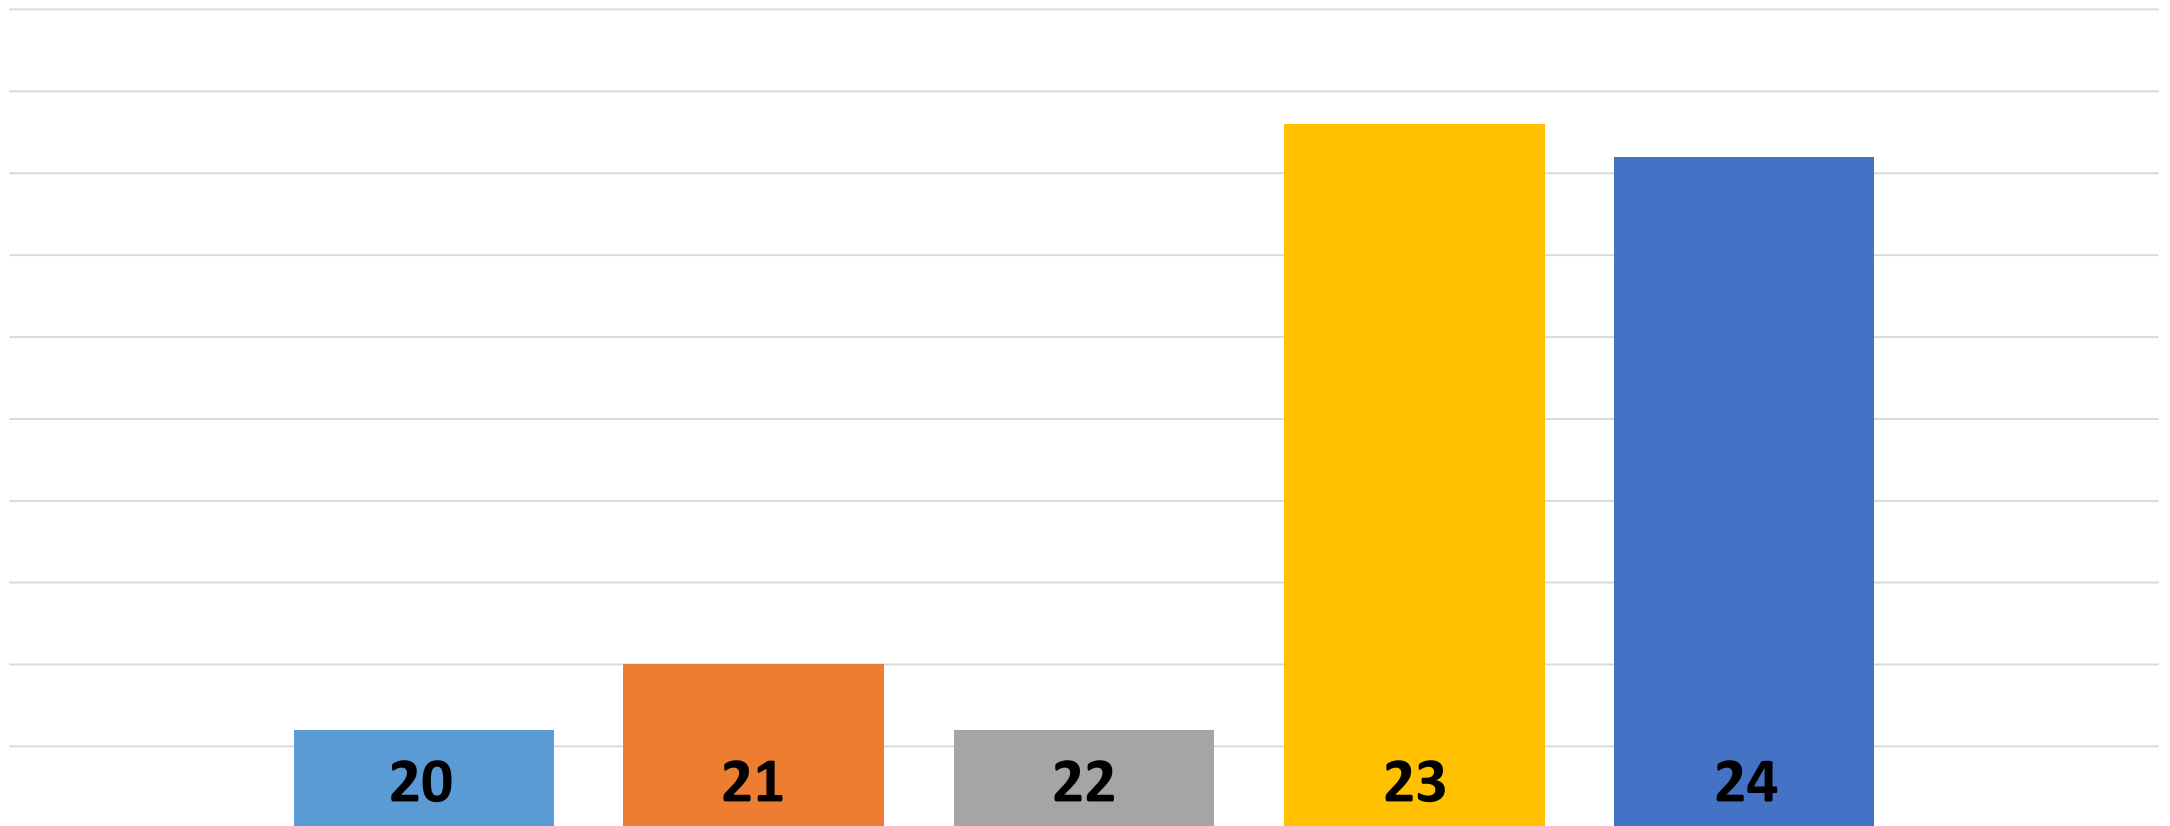
SSO

# SSO

```
std::string make_string_20() { return "12345678901234567890"; }  
std::string make_string_21() { return "123456789012345678901"; }  
std::string make_string_22() { return "1234567890123456789012"; }  
std::string make_string_23() { return "12345678901234567890123"; }  
std::string make_string_24() { return "123456789012345678901234"; }
```

# SSO

make\_string



# SSO

```
class string
{
    char* data; // 8 byte
    size_t len; // 8 byte
    size_t cap; // 8 byte
};
```

# SSO

```
class string
{
    char* data; // 8 byte
    size_t len; // 8 byte
    size_t cap; // 8 byte
};
// SSO: 23 + 1 bytes
```



# SSO

```
class string
{
    char* data; // 8 byte
    size_t len; // 8 byte
    size_t cap; // 8 byte
};
// SSO: 23 + 1 bytes
```

## Actually sizeof(std::string) == 32 ...

copy elision / RVO / NRVO.

# Antiquity

# copy elision / RVO / NRVO. Antique times

```
House build_house() {  
    House house;  
    house.add_roof(make_roof());  
    return house;  
}
```

# copy elision / RVO / NRVO. Antique times

## RVO

```
House build_house() {  
    return House(make_roof());  
}
```

# copy elision / RVO / NRVO. Antique times

## RVO

```
House build_house() {  
    return House(make_roof());  
}
```

## NRVO

```
House build_house() {  
    House house;  
    house.add_roof(make_roof());  
    return house;  
}
```

copy elision / RVO / NRVO. Antique times

NRVO fails if:

# copy elision / RVO / NRVO. Antique times

NRVO fails if:

- no single automatic storage duration object to return

```
House build_house(int id) {  
    House stone_house;  
    stone_house.add_roof(make_stone_roof());  
  
    House brick_house;  
    brick_house.add_roof(make_brick_roof());  
  
    return is_stone_house(id) ? stone_house : brick_house;  
}
```

# copy elision / RVO / NRVO. Antique times

NRVO fails if:

- no single automatic storage duration object to return
- return from function parameter

```
House build_house(House house) {  
    house.add_roof(make_roof());  
    return house;  
}
```



# copy elision / RVO / NRVO. Antique times

NRVO fails if:

- no single automatic storage duration object to return
- return from function parameter
- another return type

```
House build_house() {  
    DerivedHouse house;  
    house.add_roof(make_());  
    return house;  
}
```

copy elision / RVO / NRVO.

since C++17

# copy elision / RVO / NRVO. Since C++17

- “RVO is mandatory”

```
House build_house() {  
    return House(make_roof());  
}
```

```
House build_house(bool stone) {  
    if (stone)  
        return House(make_stone_roof());  
    else  
        return House(make_brick_roof());  
}
```

# copy elision / RVO / NRVO. Since C++17

- “RVO is mandatory”
- “unmaterialized value passing”

```
T f() {  
    return T();  
}
```

```
// only one call to default constructor of T, to initialize x  
T x = T(f());
```

*std::to\_string* reference implementation

# *std::to\_string* reference implementation

```
string to_string(int val)
{ return as_string(snprintf, initial_string<string, int>()(), "%d", val); }
```

```
string to_string(long val)
{ return as_string(snprintf, initial_string<string, long>()(), "%ld", val); }
```

...

```
wstring to_wstring(int val)
{ return as_string(get_swprintf(), initial_string<wstring, int>()(), L"%d", val); }
```

```
wstring to_wstring(long val)
{ return as_string(get_swprintf(), initial_string<wstring, long>()(), L"%ld", val); }
```

...

# *std::to\_string* reference implementation

```
string to_string(int val)
{ return as_string(snprintf, initial_string<string, int>()(), "%d", val); }
```

```
string to_string(long val)
{ return as_string(snprintf, initial_string<string, long>()(), "%ld", val); }
```

...

```
wstring to_wstring(int val)
{ return as_string(get_swprintf(), initial_string<wstring, int>()(), L"%d", val); }
```

```
wstring to_wstring(long val)
{ return as_string(get_swprintf(), initial_string<wstring, long>()(), L"%ld", val); }
```

...

# *std::to\_string* reference implementation

```
string to_string(int val)
{ return as_string(snprintf, initial_string<string, int>()(), "%d", val); }
```

```
string to_string(long val)
{ return as_string(snprintf, initial_string<string, long>()(), "%ld", val); }
```

...

```
wstring to_wstring(int val)
{ return as_string(get_swprintf(), initial_string<wstring, int>()(), L"%d", val); }
```

```
wstring to_wstring(long val)
{ return as_string(get_swprintf(), initial_string<wstring, long>()(), L"%ld", val); }
```

...



# *std::to\_string* reference implementation

```
string to_string(int val)
{ return as_string(snprintf, initial_string<string, int>(()()), "%d", val); }
```

```
string to_string(long val)
{ return as_string(snprintf, initial_string<string, long>(()()), "%ld", val); }
```

...

```
wstring to_wstring(int val)
{ return as_string(get_swprintf(), initial_string<wstring, int>(()()), L"%d", val); }
```

```
wstring to_wstring(long val)
{ return as_string(get_swprintf(), initial_string<wstring, long>(()()), L"%ld", val); }
```

...

# *std::to\_string* reference implementation

```
string to_string(int val)
{ return as_string(snprintf, initial_string<string, int>()( ), "%d", val); }
```

```
string to_string(long val)
{ return as_string(snprintf, initial_string<string, long>()( ), "%ld", val); }
```

...

```
wstring to_wstring(int val)
{ return as_string(get_swprintf(), initial_string<wstring, int>()( ), L"%d", val); }
```

```
wstring to_wstring(long val)
{ return as_string(get_swprintf(), initial_string<wstring, long>()( ), L"%ld", val); }
```

...

# *std::to\_string* reference implementation

```
string to_string(int val)
{ return as_string(snprintf, initial_string<string, int>()( ), "%d", val); }
```

```
string to_string(long val)
{ return as_string(snprintf, initial_string<string, long>()( ), "%ld", val); }
```

...

```
wstring to_wstring(int val)
{ return as_string(get_swprintf(), initial_string<wstring, int>()( ), L"%d", val); }
```

```
wstring to_wstring(long val)
{ return as_string(get_swprintf(), initial_string<wstring, long>()( ), L"%ld", val); }
```

...

# *std::to\_string* reference implementation

```
template <class S, class V>  
struct initial_string;
```

# *std::to\_string* reference implementation

```
template <class S, class V>  
struct initial_string;
```

# *std::to\_string* reference implementation

```
template <class S, class V>  
struct initial_string;
```

# *std::to\_string* reference implementation

```
template <class S, class V>  
struct initial_string;
```

```
template <class V>  
struct initial_string<string, V>  
{  
    string operator()() const  
    {  
        string s;  
        s.resize(s.capacity());  
        return s;  
    }  
};
```

# *std::to\_string* reference implementation

```
template <class S, class V>  
struct initial_string;
```

```
template <class V>  
struct initial_string<string, V>  
{  
    string operator()() const  
    {  
        string s;  
        s.resize(s.capacity());  
        return s;  
    }  
};
```



# *std::to\_string* reference implementation

```
template <class S, class V>  
struct initial_string;
```

```
template <class V>  
struct initial_string<string, V>  
{  
    string operator()() const  
    {  
        string s;  
        s.resize(s.capacity());  
        return s;  
    }  
};
```

# *std::to\_string* reference implementation

```
template <class S, class V>
struct initial_string;
```

```
template <class V>
struct initial_string<string, V>
{
    string operator()() const
    {
        string s;
        s.resize(s.capacity());
        return s;
    }
};
```

```
template <class V>
struct initial_string<wstring, V>
{
    wstring operator()() const
    {
        wstring s(23, wchar_t()); // 23 for ull
        s.resize(s.capacity());
        return s;
    }
};
```

# *std::to\_string* reference implementation

```
template <class S, class V>
struct initial_string;
```

```
template <class V>
struct initial_string<string, V>
{
    string operator()() const
    {
        string s;
        s.resize(s.capacity());
        return s;
    }
};
```

```
template <class V>
struct initial_string<wstring, V>
{
    wstring operator()() const
    {
        wstring s(23, wchar_t()); // 23 for ull
        s.resize(s.capacity());
        return s;
    }
};
```

# *std::to\_string* reference implementation

```
template <class S, class V>
struct initial_string;
```

```
template <class V>
struct initial_string<string, V>
{
    string operator()() const
    {
        string s;
        s.resize(s.capacity());
        return s;
    }
};
```

```
template <class V>
struct initial_string<wstring, V>
{
    wstring operator()() const
    {
        wstring s(23, wchar_t()); // 23 for ull
        s.resize(s.capacity());
        return s;
    }
};
```

# *std::to\_string* reference implementation

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, S s, const typename S::value_type* fmt, V a) {
    size_t available = s.size();
    while (true) {
        int status = sprintf_like(&s[0], available + 1, fmt, a);
        if (status >= 0) {
            size_t used = static_cast<size_t>(status);
            if (used <= available) {
                s.resize(used); // success: fit size
                break;          // success: return
            }
            available = used; // assume this is advice of how much space we need.
        } else {
            available = available * 2 + 1;
        }
        s.resize(available);
    }
    return s;
}
```

# *std::to\_string* reference implementation

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, S s, const typename S::value_type* fmt, V a) {
    size_t available = s.size();
    while (true) {
        int status = sprintf_like(&s[0], available + 1, fmt, a);
        if (status >= 0) {
            size_t used = static_cast<size_t>(status);
            if (used <= available) {
                s.resize(used); // success: fit size
                break;          // success: return
            }
            available = used; // assume this is advice of how much space we need.
        } else {
            available = available * 2 + 1;
        }
        s.resize(available);
    }
    return s;
}
```

# *std::to\_string* reference implementation

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, S s, const typename S::value_type* fmt, V a) {
    size_t available = s.size();
    while (true) {
        int status = sprintf_like(&s[0], available + 1, fmt, a);
        if (status >= 0) {
            size_t used = static_cast<size_t>(status);
            if (used <= available) {
                s.resize(used); // success: fit size
                break;          // success: return
            }
            available = used; // assume this is advice of how much space we need.
        } else {
            available = available * 2 + 1;
        }
        s.resize(available);
    }
    return s;
}
```

# *std::to\_string* reference implementation

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, S s, const typename S::value_type* fmt, V a) {
    size_t available = s.size();
    while (true) {
        int status = sprintf_like(&s[0], available + 1, fmt, a);
        if (status >= 0) {
            size_t used = static_cast<size_t>(status);
            if (used <= available) {
                s.resize(used); // success: fit size
                break;          // success: return
            }
            available = used; // assume this is advice of how much space we need.
        } else {
            available = available * 2 + 1;
        }
        s.resize(available);
    }
    return s;
}
```



# *std::to\_string* reference implementation

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, S s, const typename S::value_type* fmt, V a) {
    size_t available = s.size();
    while (true) {
        int status = sprintf_like(&s[0], available + 1, fmt, a);
        if (status >= 0) {
            size_t used = static_cast<size_t>(status);
            if (used <= available) {
                s.resize(used); // success: fit size
                break;          // success: return
            }
            available = used; // assume this is advice of how much space we need.
        } else {
            available = available * 2 + 1;
        }
        s.resize(available);
    }
    return s;
}
```

# *std::to\_string* reference implementation

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, S s, const typename S::value_type* fmt, V a) {
    size_t available = s.size();
    while (true) {
        int status = sprintf_like(&s[0], available + 1, fmt, a);
        if (status >= 0) {
            size_t used = static_cast<size_t>(status);
            if (used <= available) {
                s.resize(used); // success: fit size
                break;          // success: return
            }
            available = used; // assume this is advice of how much space we need.
        } else {
            available = available * 2 + 1;
        }
        s.resize(available);
    }
    return s;
}
```

# *std::to\_string* reference implementation

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, S s, const typename S::value_type* fmt, V a) {
    size_t available = s.size();
    while (true) {
        int status = sprintf_like(&s[0], available + 1, fmt, a);
        if (status >= 0) {
            size_t used = static_cast<size_t>(status);
            if (used <= available) {
                s.resize(used); // success: fit size
                break;          // success: return
            }
            available = used; // assume this is advice of how much space we need.
        } else {
            available = available * 2 + 1;
        }
        s.resize(available);
    }
    return s;
}
```

# *std::to\_string* reference implementation

- *to\_string*:
  - SSO + *sprintf*
- *to\_wstring*:
  - allocation + *sprintf* (missing SSO opportunities)
- missing copy elision opportunities

# *std::to\_string* reference implementation

- *to\_string*:
  - SSO + *sprintf*
- *to\_wstring*:
  - allocation + *sprintf* (missing SSO opportunities)
- missing copy elision opportunities

Test matrix	1	ULLONG_MAX
<i>to_string</i>	<i>to_string(1)</i>	<i>to_string(ULLONG_MAX)</i>
<i>to_wstring</i>	<i>to_wstring(1)</i>	<i>to_wstring(ULLONG_MAX)</i>

# Proposal 1

# Proposal 1

Idea:

- use char buffer on stack
- call sprintf into stack buffer
- return string from buffer on success
- fallback to the previous algorithm on failure

# Proposal 1

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, const typename S::value_type* fmt, V a)
{
    // fast path for nice sprintf functions
    constexpr size_t size = BIG_ENOUGH_SIZE_FOR_TYPE_V;
    typename S::value_type tmp[size] = {};
    const int len = sprintf_like(tmp, size, fmt, a);
    if (len <= size)
        return S(tmp, tmp + len); // copy elision guarantee since C++17

    // fallback to previous algorithm for weird sprintf functions
    S s;
    ...
    return s;
}
```



# Proposal 1

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, const typename S::value_type* fmt, V a)
{
    // fast path for nice sprintf functions
    constexpr size_t size = BIG_ENOUGH_SIZE_FOR_TYPE_V;
    typename S::value_type tmp[size] = {};
    const int len = sprintf_like(tmp, size, fmt, a);
    if (len <= size)
        return S(tmp, tmp + len); // copy elision guarantee since C++17

    // fallback to previous algorithm for weird sprintf functions
    S s;
    ...
    return s;
}
```

# Proposal 1

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, const typename S::value_type* fmt, V a)
{
    // fast path for nice sprintf functions
    constexpr size_t size = BIG_ENOUGH_SIZE_FOR_TYPE_V;
    typename S::value_type tmp[size] = {};
    const int len = sprintf_like(tmp, size, fmt, a);
    if (len <= size)
        return S(tmp, tmp + len); // copy elision guarantee since C++17

    // fallback to previous algorithm for weird sprintf functions
    S s;
    ...
    return s;
}
```

# Proposal 1

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, const typename S::value_type* fmt, V a)
{
    // fast path for nice sprintf functions
    constexpr size_t size = BIG_ENOUGH_SIZE_FOR_TYPE_V;
    typename S::value_type tmp[size] = {};
    const int len = sprintf_like(tmp, size, fmt, a);
    if (len <= size)
        return S(tmp, tmp + len); // copy elision guarantee since C++17

    // fallback to previous algorithm for weird sprintf functions
    S s;
    ...
    return s;
}
```

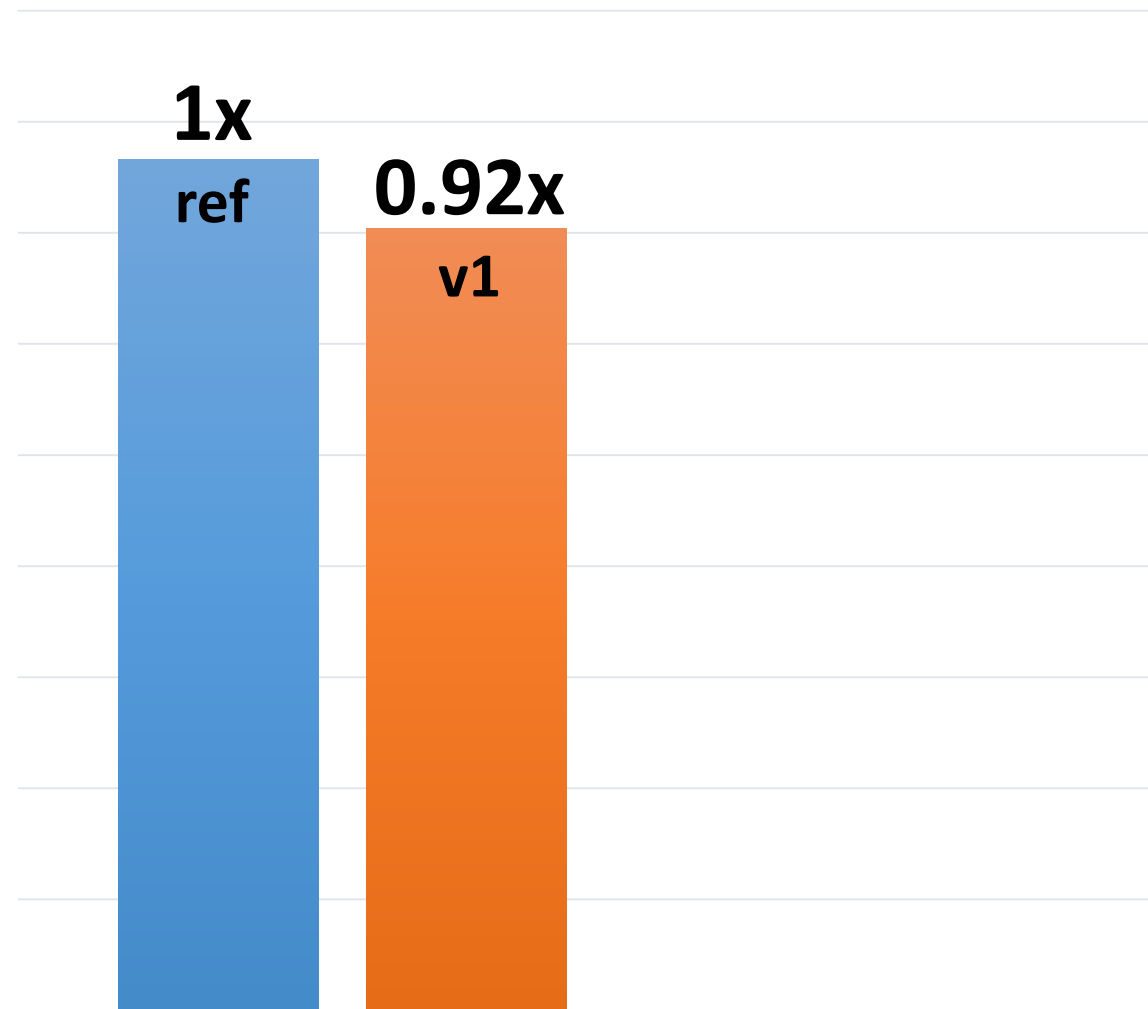
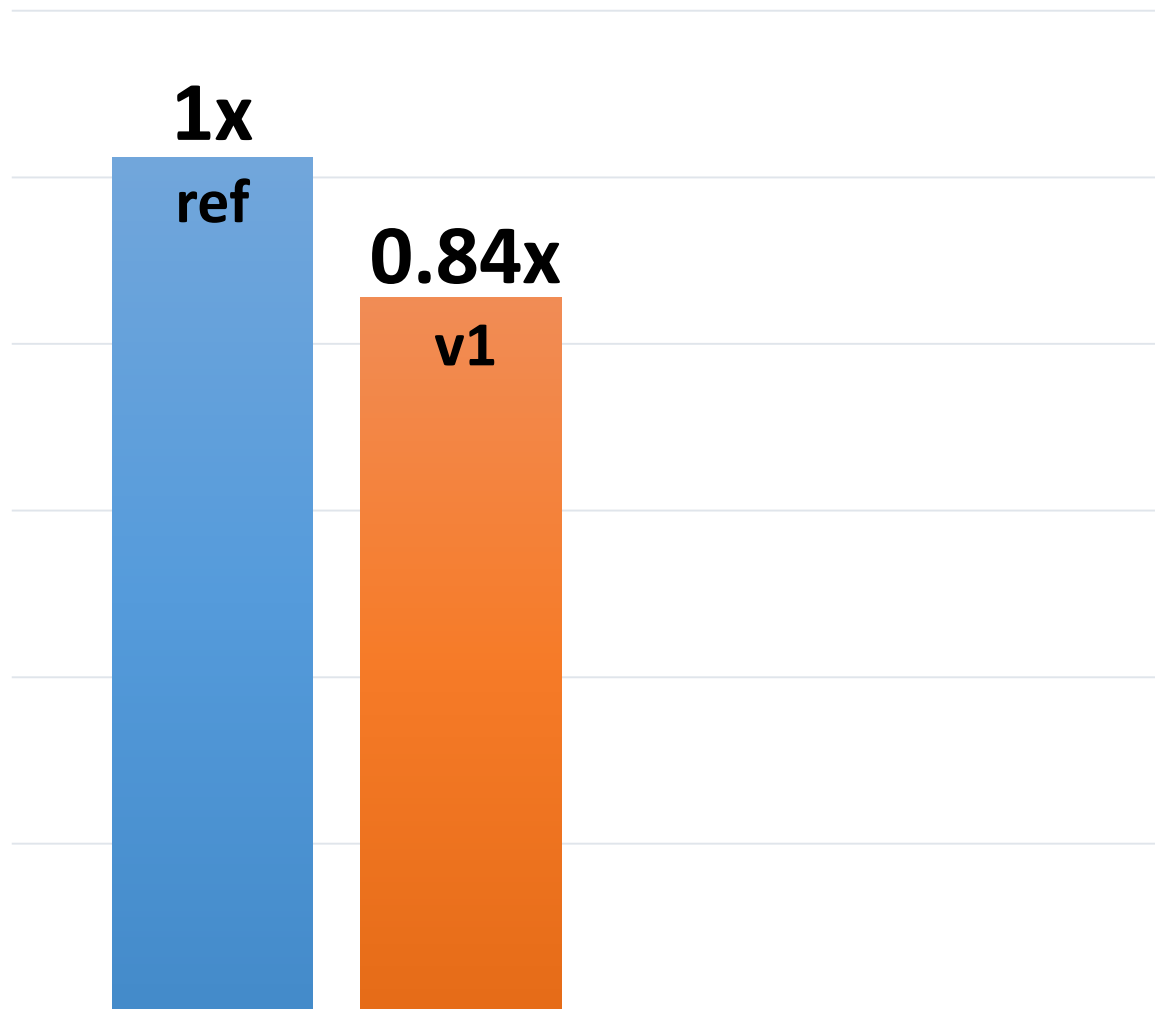
# Proposal 1

```
template<typename S, typename P, typename V>
inline S as_string(P sprintf_like, const typename S::value_type* fmt, V a)
{
    // fast path for nice sprintf functions
    constexpr size_t size = BIG_ENOUGH_SIZE_FOR_TYPE_V;
    typename S::value_type tmp[size] = {};
    const int len = sprintf_like(tmp, size, fmt, a);
    if (len <= size)
        return S(tmp, tmp + len); // copy elision guarantee since C++17

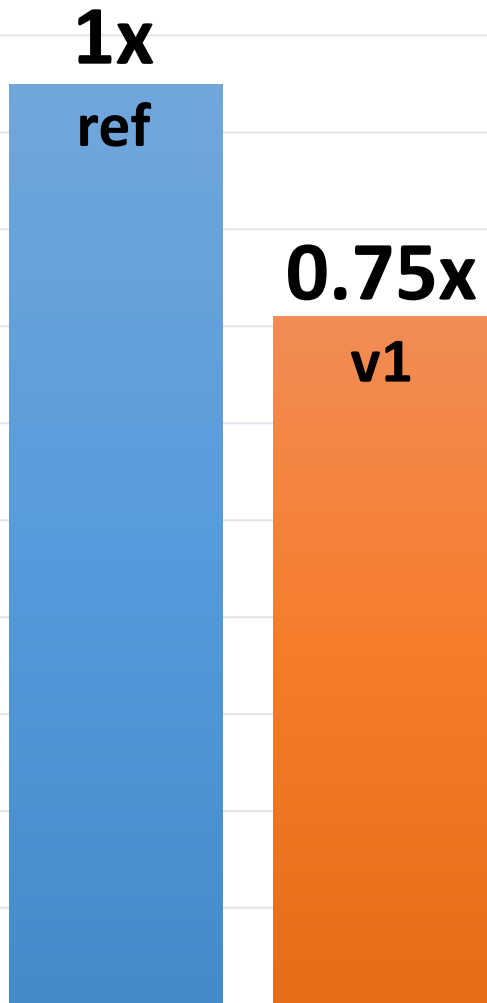
    // fallback to previous algorithm for weird sprintf functions
    S s;
    ...
    return s;
}
```

*to\_string(1)*

*to\_string(max)*



*to\_wstring(1)*



*to\_wstring(max)*



25% speedup... seems nice?





# Proposal 2

# Proposal 2

Ed Schouten: “If performance matters, it makes sense to handroll this for the integer cases”

# Proposal 2

- use *val % 10* and *val / 10* (2 division ops per char)

# Proposal 2

- use *val % 10* and *val / 10* (2 division ops per char)
- fill stack buffer from right to left

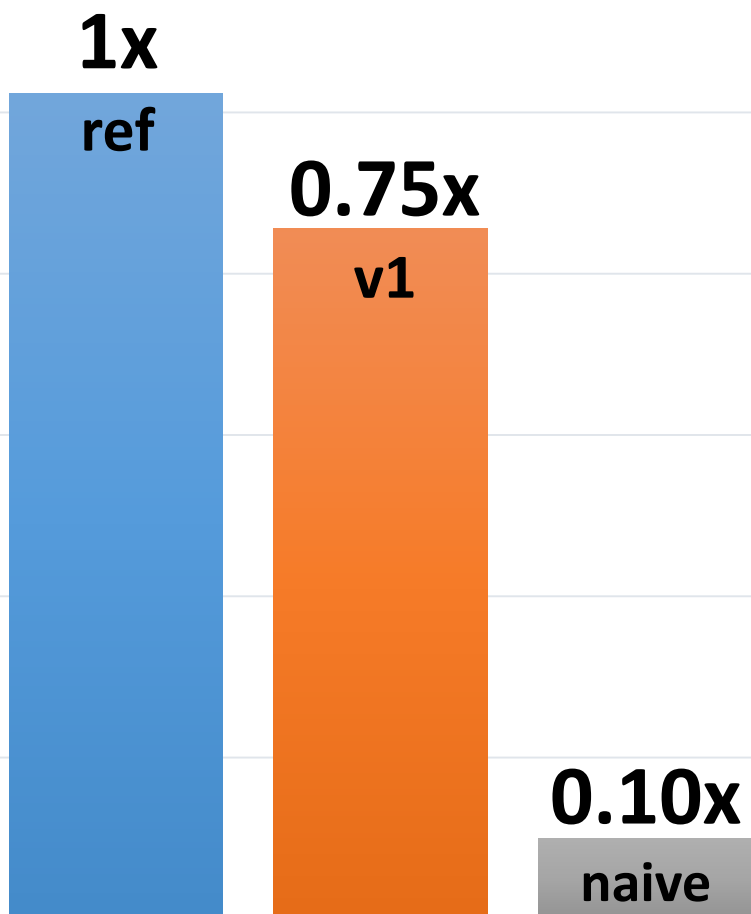
# Proposal 2

- use *val % 10* and *val / 10* (2 division ops per char)
- fill stack buffer from right to left
- do not forget about negative numbers

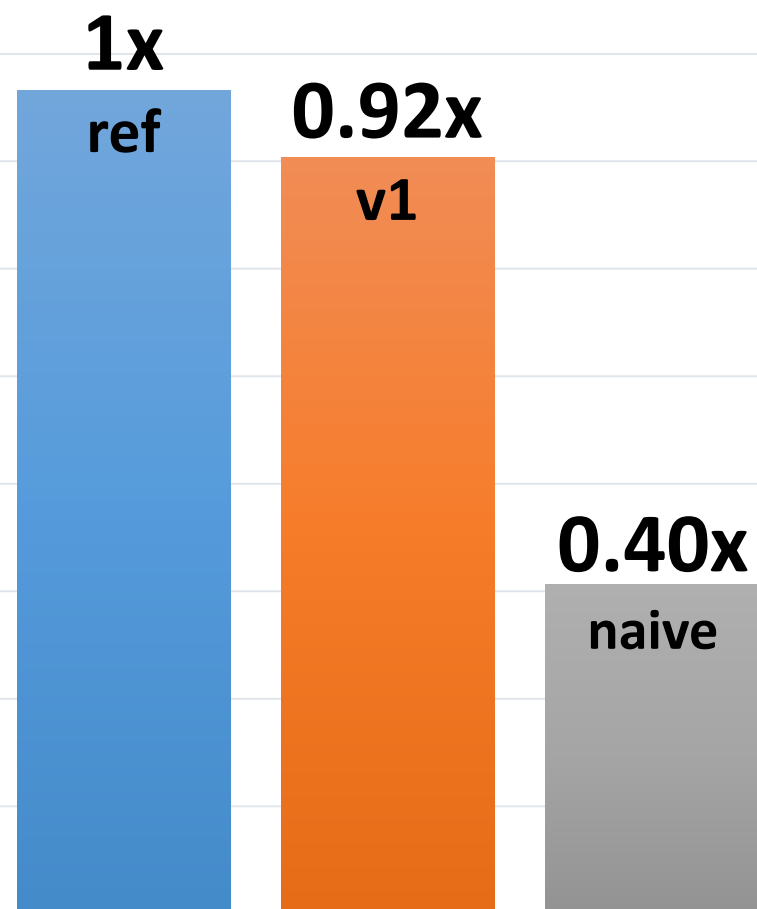
# Proposal 2

- use *val % 10* and *val / 10* (2 division ops per char)
- fill stack buffer from right to left
- do not forget about negative numbers
  - $(val \% 10) \leq 0$  according to C++ standard par. "Multiplicative operators" (ref. to "truncation towards zero" since C99)
  - need to deal with sign
  - `std::numeric_limits<V>::min()`

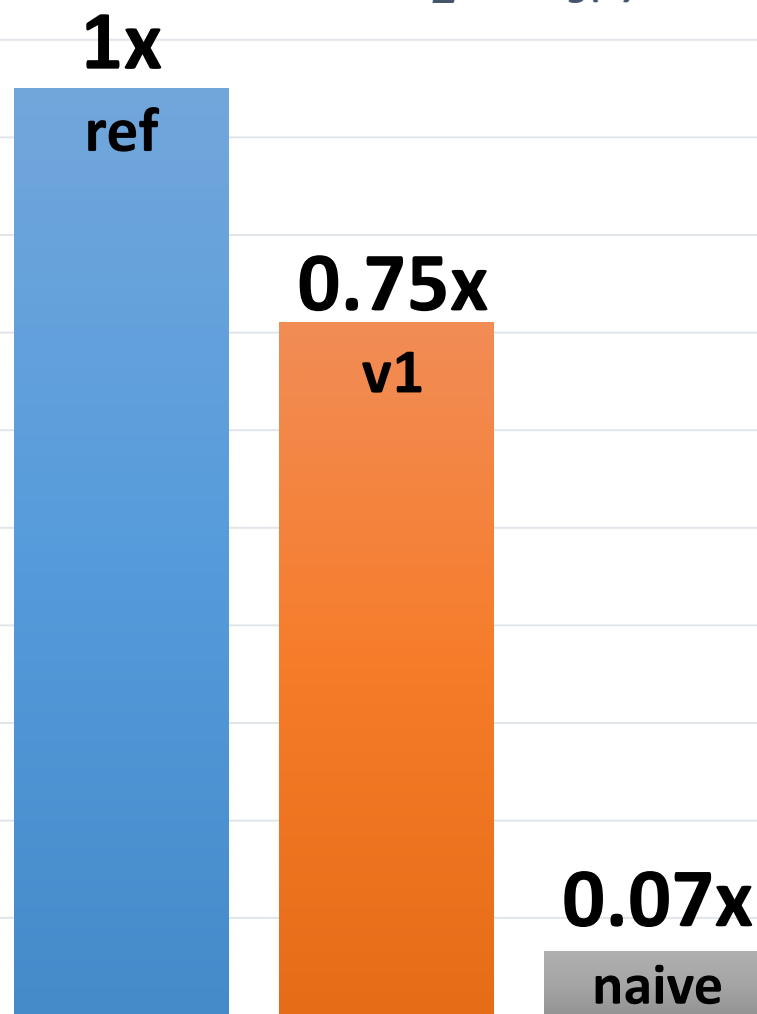
*to\_string(1)*



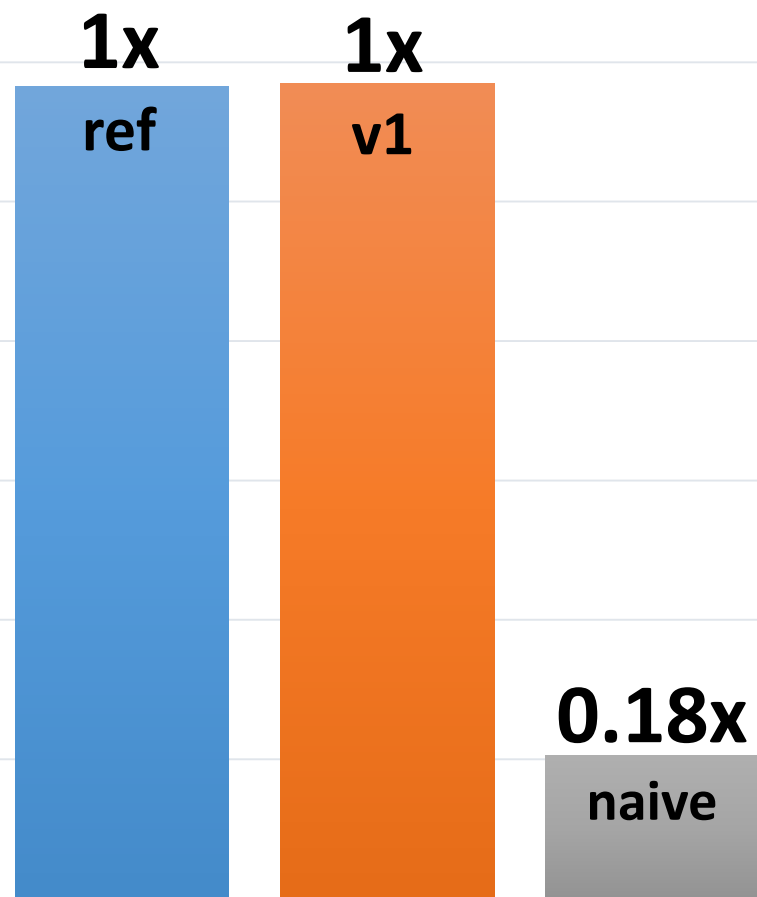
*to\_string(max)*



*to\_wstring(1)*



*to\_wstring(max)*





# Why sprintf is so slow?

?

# Why sprintf is so slow?

- it is general-purpose algorithm

# Why sprintf is so slow?

- it is general-purpose algorithm
- format parsing?

# Why sprintf is so slow?

- it is general-purpose algorithm
- format parsing?
- varargs?

# Why sprintf is so slow?

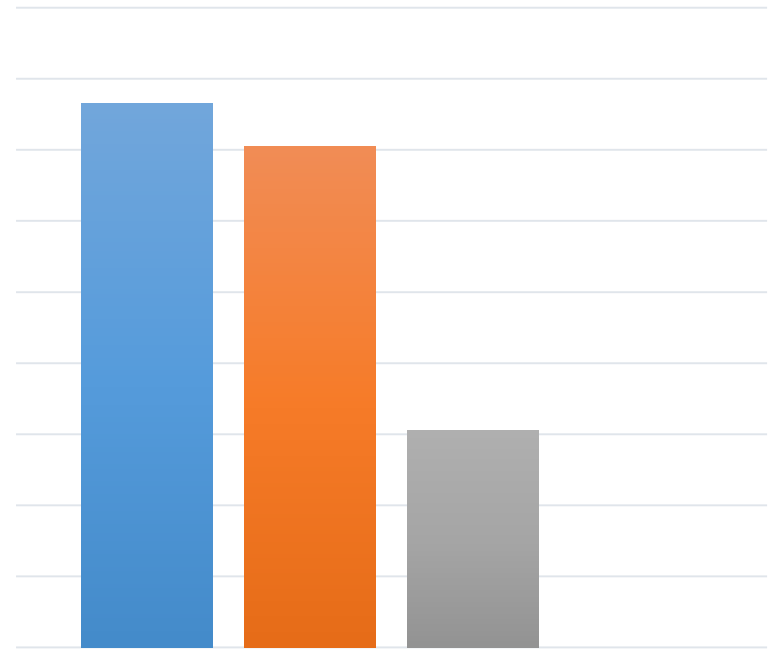
- it is general-purpose algorithm
- format parsing?
- varargs?
- locale and synchronizations?

# Why sprintf is so slow?

- it is general-purpose algorithm
- format parsing?
- varargs?
- locale and synchronizations?

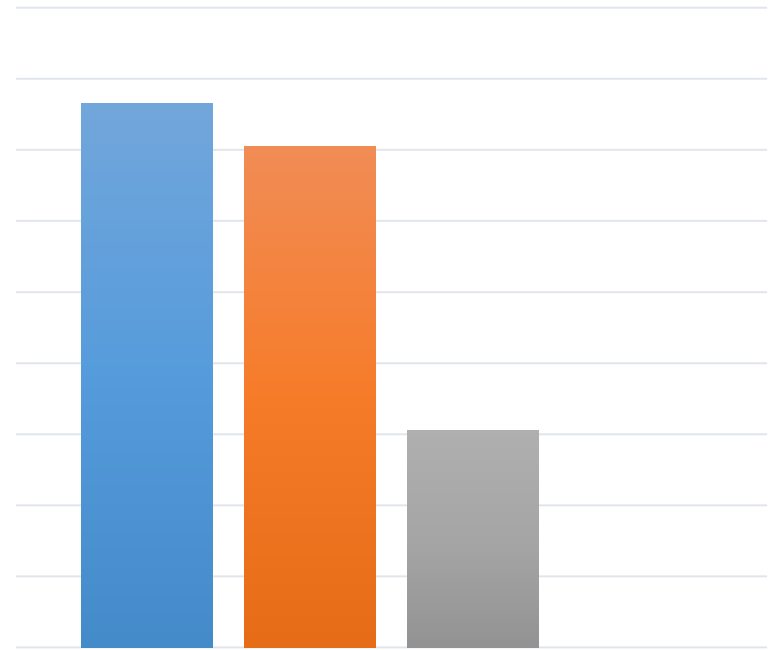
Let's profile it!

# sprintf profile for ULLONG\_MAX



%	cumulative	self		self	total		
time	seconds	seconds	calls	Ts /	call	Ts /	call
38.99	43.04	43.04					_itoa_word
29.43	75.53	32.49					vfprintf
10.84	87.50	11.97					_IO_default_xsputn
4.22	92.16	4.66					__vsprintf_chk
4.16	96.75	4.60					__strchrnul_avx2
3.18	100.27	3.52					_IO_str_init_static_internal
2.60	103.13	2.87					_IO_no_init
...							

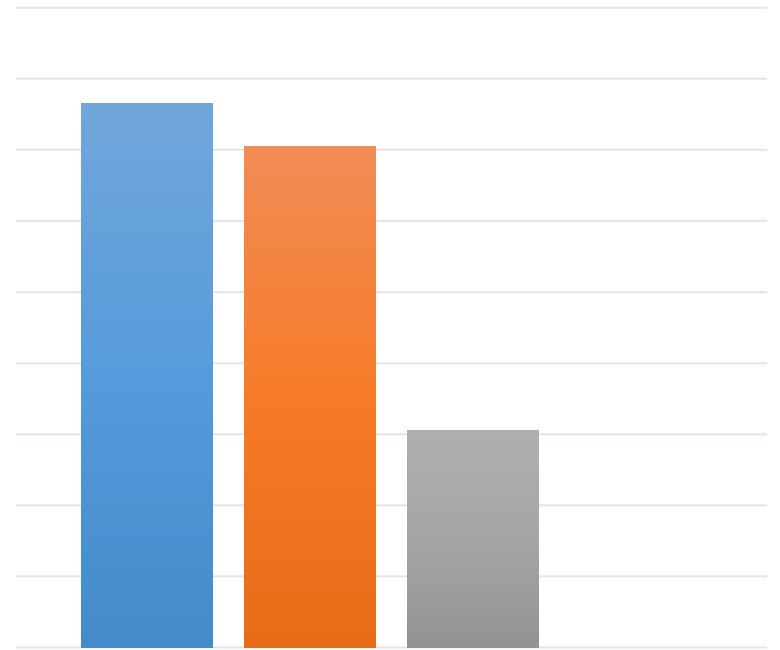
# sprintf profile for ULLONG\_MAX



% time	cumulative seconds	self seconds	calls	self Ts / call	total Ts / call	name
38.99	43.04	43.04				_itoa_word
29.43	75.53	32.49				vfprintf
10.84	87.50	11.97				_IO_default_xsputn
4.22	92.16	4.66				__vsprintf_chk
4.16	96.75	4.60				__strchrnul_avx2
3.18	100.27	3.52				_IO_str_init_static_internal
2.60	103.13	2.87				_IO_no_init
...						

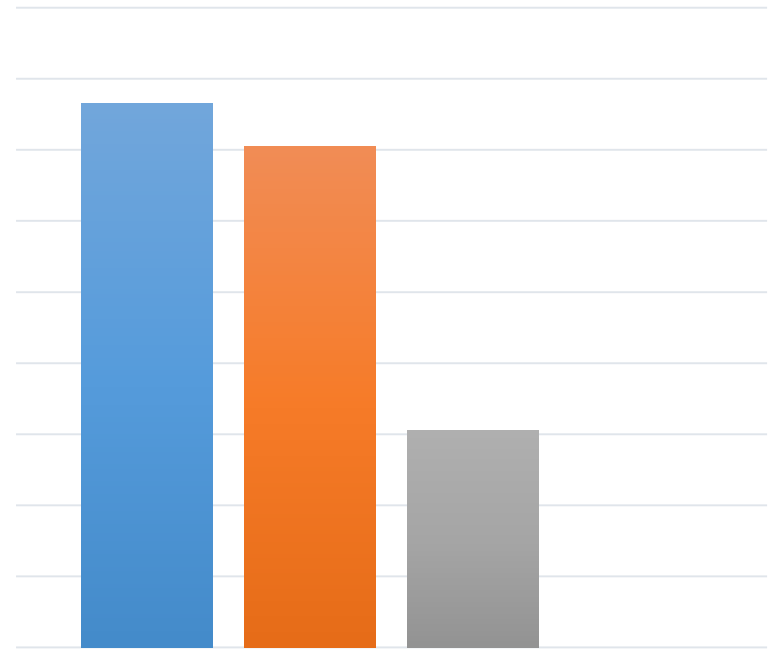


# sprintf profile for ULLONG\_MAX



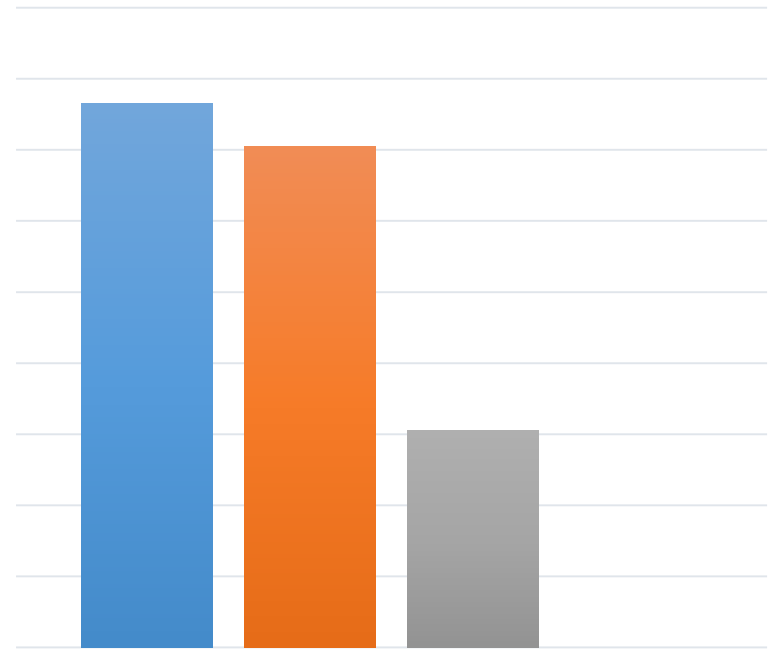
%	cumulative	self		self	total		name
time	seconds	seconds	calls	Ts /	call	Ts /	call
38.99	43.04	43.04					_itoa_word
29.43	75.53	32.49					vfprintf
10.84	87.50	11.97					_IO_default_xspn
4.22	92.16	4.66					__vsprintf_chk
4.16	96.75	4.60					__strchrnul_avx2
3.18	100.27	3.52					_IO_str_init_static_internal
2.60	103.13	2.87					_IO_no_init
...							

# sprintf profile for ULLONG\_MAX



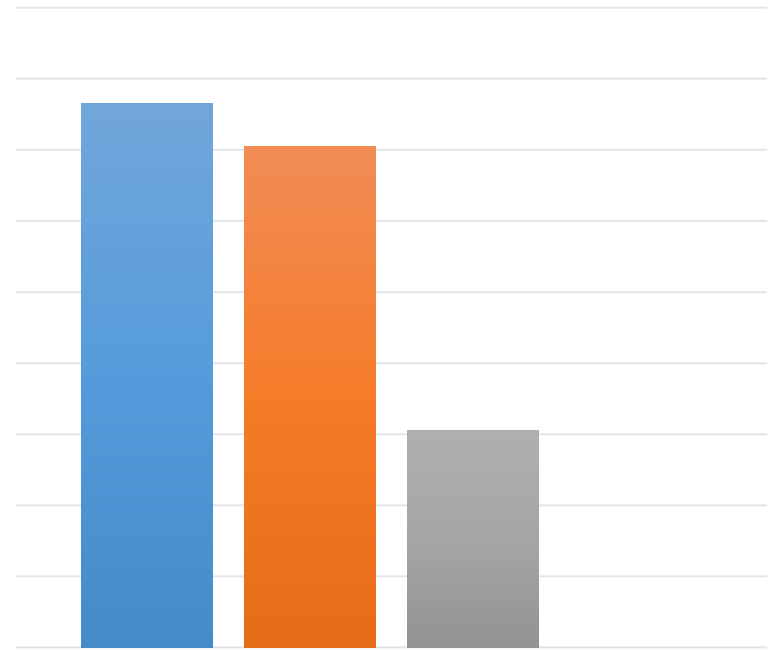
%	cumulative	self		self	total		name
time	seconds	seconds	calls	Ts /	call	Ts /	call
38.99	43.04	43.04					_itoa_word
29.43	75.53	32.49					vfprintf
10.84	87.50	11.97					_IO_default_xsputn
4.22	92.16	4.66					__vsprintf_chk
4.16	96.75	4.60					__strchrnul_avx2
3.18	100.27	3.52					_IO_str_init_static_internal
2.60	103.13	2.87					_IO_no_init
...							

# sprintf profile for ULLONG\_MAX



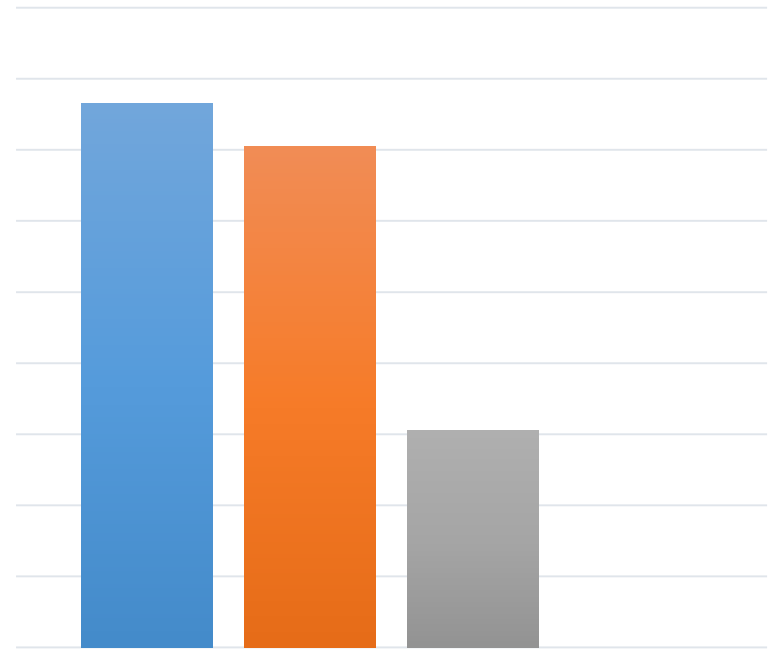
%	cumulative	self		self	total	
time	seconds	seconds	calls	Ts /	Ts /	call name
38.99	43.04	43.04		call	call	_itoa_word
29.43	75.53	32.49				vfprintf
10.84	87.50	11.97				_IO_default_xspn
4.22	92.16	4.66				__vsprintf_chk
4.16	96.75	4.60				__strchrnul_avx2
3.18	100.27	3.52				_IO_str_init_static_internal
2.60	103.13	2.87				_IO_no_init
...						

# sprintf profile for ULLONG\_MAX



%	cumulative	self		self	total	
time	seconds	seconds	calls	Ts / call	Ts / call	name
38.99	43.04	43.04				_itoa_word
29.43	75.53	32.49				vfprintf
10.84	87.50	11.97				_IO_default_xsputn
4.22	92.16	4.66				__vsprintf_chk
4.16	96.75	4.60				__strchrnul_avx2
3.18	100.27	3.52				_IO_str_init_static_internal
2.60	103.13	2.87				_IO_no_init
...						

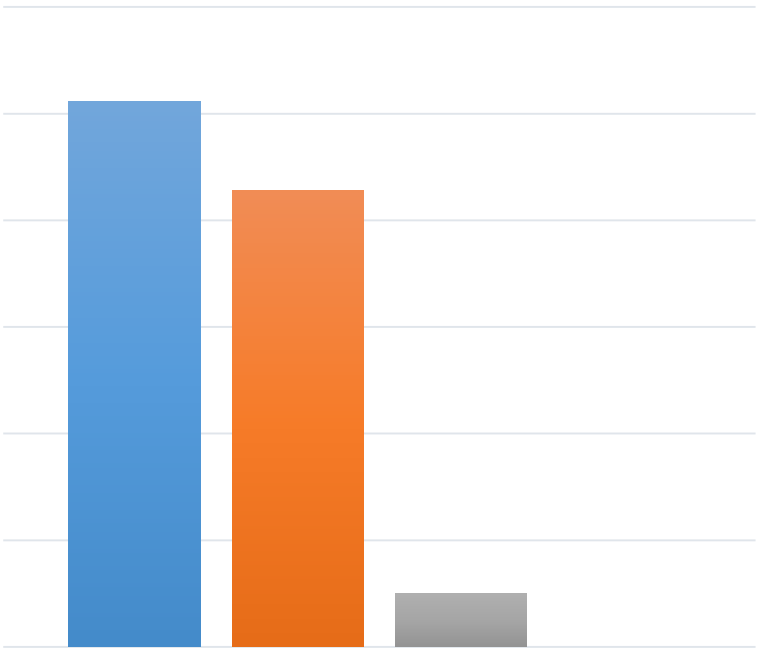
# sprintf profile for ULLONG\_MAX



%	cumulative	self		self	total		
time	seconds	seconds	calls	Ts /	call	Ts /	call
38.99	43.04	43.04					_itoa_word
29.43	75.53	32.49					vfprintf
10.84	87.50	11.97					_IO_default_xsputn
4.22	92.16	4.66					__vsprintf_chk
4.16	96.75	4.60					__strchrnul_avx2
3.18	100.27	3.52					_IO_str_init_static_internal
2.60	103.13	2.87					_IO_no_init

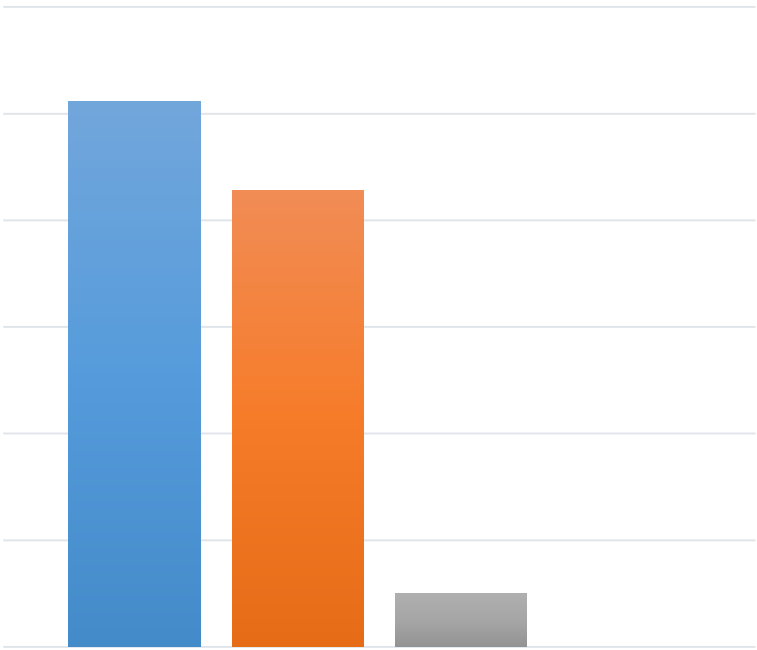
...

# sprintf profile for 1



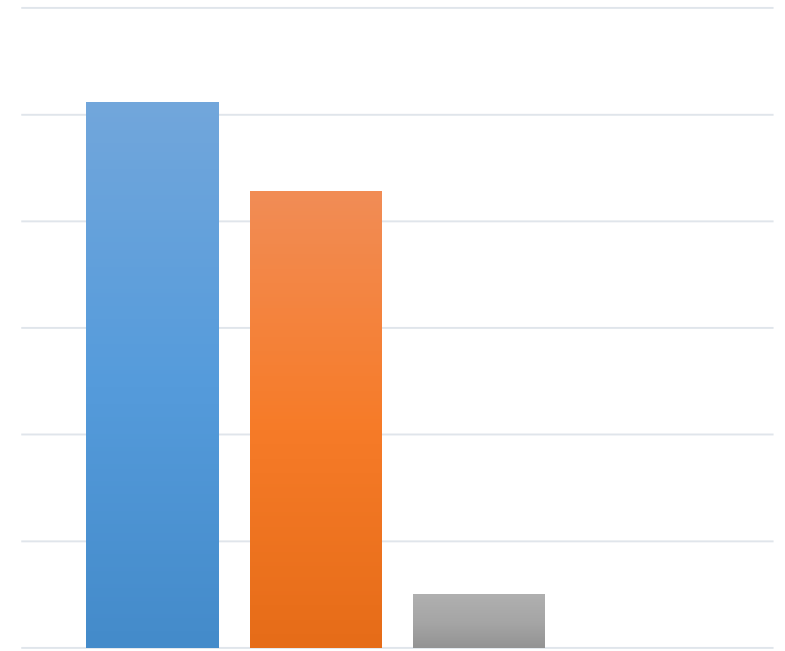
%	cumulative	self		self	total		
time	seconds	seconds	calls	Ts /	call	Ts /	call
47.45	31.87	31.87					name
							vfprintf
10.35	38.82	6.95					__strchrnul_avx2
8.14	44.29	5.47					_IO_default_xsputn
6.97	48.97	4.68					__vsprintf_chk
...							

# sprintf profile for 1



%	cumulative	self		self	total		
time	seconds	seconds	calls	Ts /	call	Ts /	call
47.45	31.87	31.87					name
10.35	38.82	6.95					vfprintf
8.14	44.29	5.47					__strchrnul_avx2
6.97	48.97	4.68					_IO_default_xsputn
...							__vsprintf_chk

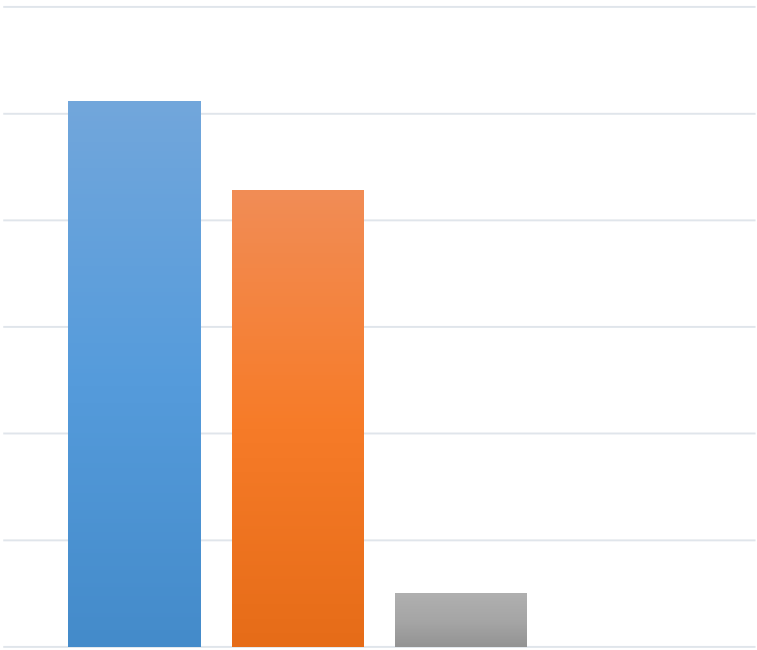
# sprintf profile for 1



%	cumulative	self		self	total		name
time	seconds	seconds	calls	Ts /	call	Ts /	call
47.45	31.87	31.87					vfprintf
10.35	38.82	6.95					__strchrnul_avx2
8.14	44.29	5.47					_IO_default_xsputn
6.97	48.97	4.68					__vsprintf_chk
...							



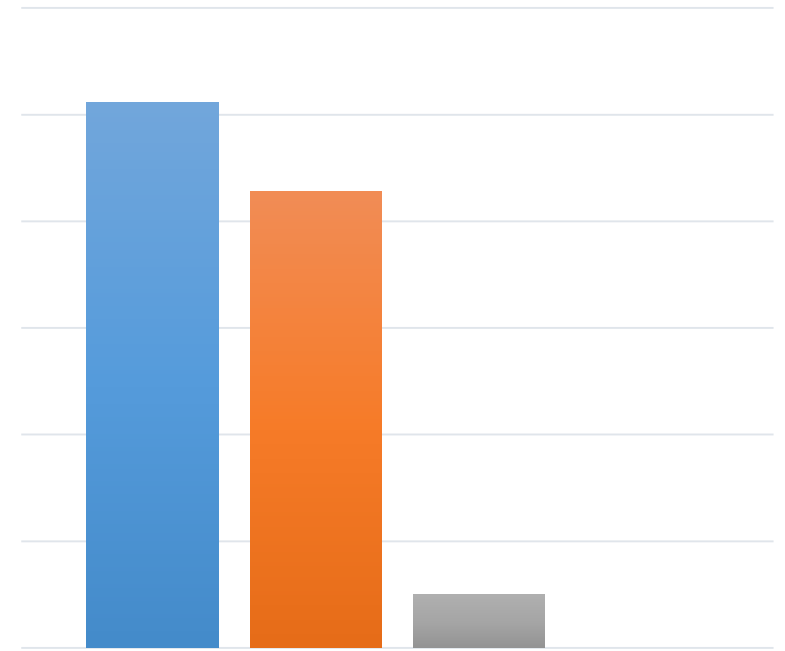
# sprintf profile for 1



%	cumulative	self		self	total		
time	seconds	seconds	calls	Ts /	call	Ts /	call
47.45	31.87	31.87					vfprintf
10.35	38.82	6.95					__strchrnul_avx2
8.14	44.29	5.47					_IO_default_xsputn
6.97	48.97	4.68					__vsprintf_chk

...

# sprintf profile for 1



%	cumulative	self		self	total		name
time	seconds	seconds	calls	Ts /	call	Ts /	call
47.45	31.87	31.87					vfprintf
10.35	38.82	6.95					__strchrnul_avx2
8.14	44.29	5.47					_IO_default_xsputn
6.97	48.97	4.68					__vsprintf_chk
6.70	53.47	4.50					_IO_no_init
5.07	56.88	3.41					_IO_str_init_static_internal
4.23	59.72	2.84					_itoa_word

...

sprintf profile conclusions

# sprintf profile conclusions

- format parsing is expensive

# sprintf profile conclusions

- format parsing is expensive
- internal structures overhead

# sprintf profile conclusions

- format parsing is expensive
- internal structures overhead
- sanity checks overhead

# sprintf profile conclusions

- format parsing is expensive
- internal structures overhead
- sanity checks overhead
- locale access is lazy (no synchronizations)

93% speedup... seems nice?





# Proposal 3

# Proposal 3

Marshall Clow: “just reuse *std::to\_chars*”

## std::to\_chars

---

Defined in header `<charconv>`

---

```
std::to_chars_result to_chars(char* first, char* last,  
                             /*see below*/ value, int base = 10);
```

---

(1)

```
struct to_chars_result {  
    char* ptr;  
    std::errc ec;  
};
```

---

(5)

# Proposal 3

Marshall Clow: “just reuse *std::to\_chars*”

## std::to\_chars

---

Defined in header `<charconv>`

---

```
std::to_chars_result to_chars(char* first, char* last,  
                             /*see below*/ value, int base = 10);
```

---

(1) (since C++17)

```
struct to_chars_result {  
    char* ptr;  
    std::errc ec;  
};
```

---

(5) (since C++17)

# Proposal 3

Marshall Clow: “just reuse *std::to\_chars*”

## std::to\_chars

Defined in header `<charconv>`

```
std::to_chars_result to_chars(char* first, char* last,  
                             /*see below*/ value, int base = 10);
```

(1) (since C++17)

```
struct to_chars_result {  
    char* ptr;  
    std::errc ec;  
};
```

(5) (since C++17)

## std::to\_string

Defined in header `<string>`

```
std::string to_string(int value);
```

(1) (since C++11)

# Proposal 3

Marshall Clow: “just reuse *std::to\_chars*”

Marshall Clow: “not a problem, *std::to\_chars* is since C++11 now:” (libc++ only)

<https://reviews.llvm.org/D59598>

`std::to_string`

Defined in header `<string>`

`std::string to_string( int value );`

(1) (since C++11)

(1) (since C++17)

(5) (since C++17)

# Proposal 3

```
template <typename S, typename V>
S i_to_string(const V v)
{
    constexpr size_t bufsize = numeric_limits<V>::digits10 + 2;
    char buf[bufsize];
    const auto res = to_chars(buf, buf + bufsize, v);
    return S(buf, res.ptr);
}
```

# Proposal 3

```
template <typename S, typename V>
S i_to_string(const V v)
{
    constexpr size_t bufsize = numeric_limits<V>::digits10 + 2;
    char buf[bufsize];
    const auto res = to_chars(buf, buf + bufsize, v);
    return S(buf, res.ptr);
}
```



# Proposal 3

```
template <typename S, typename V>
S i_to_string(const V v)
{
    constexpr size_t bufsize = numeric_limits<V>::digits10 + 2;
    char buf[bufsize];
    const auto res = to_chars(buf, buf + bufsize, v);
    return S(buf, res.ptr);
}
```

# Proposal 3

```
template <typename S, typename V>
S i_to_string(const V v)
{
    constexpr size_t bufsize = numeric_limits<V>::digits10 + 2;
    char buf[bufsize];
    const auto res = to_chars(buf, buf + bufsize, v);
    return S(buf, res.ptr);
}
```

# Proposal 3

```
template <typename S, typename V>
S i_to_string(const V v)
{
    constexpr size_t bufsize = numeric_limits<V>::digits10 + 2;
    char buf[bufsize];
    const auto res = to_chars(buf, buf + bufsize, v);
    return S(buf, res.ptr);
}
```

# Proposal 3

```
template <typename S, typename V>
S i_to_string(const V v)
{
    constexpr size_t bufsize = numeric_limits<V>::digits10 + 2;
    char buf[bufsize];
    const auto res = to_chars(buf, buf + bufsize, v);
    return S(buf, res.ptr);
}
```

The value of `std::numeric_limits<T>::digits10` is the number of base-10 digits that can be represented by the type `T` without change.

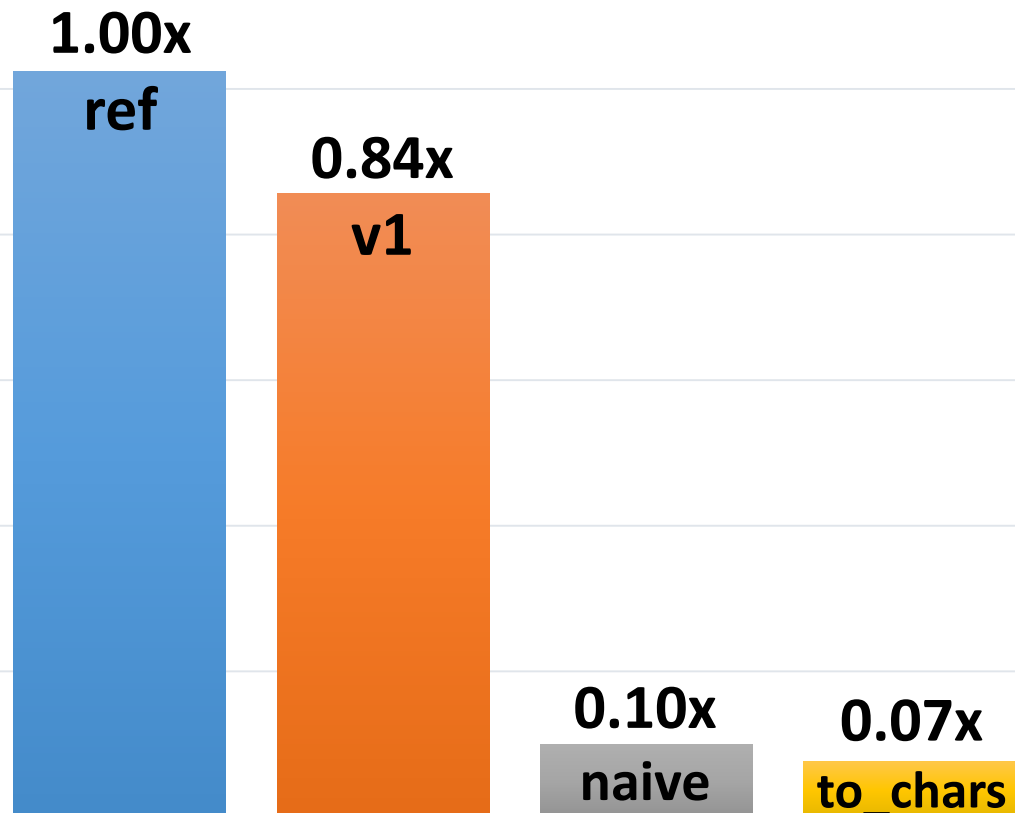
# Proposal 3

```
template <typename S, typename V>
S i_to_string(const V v)
{
    constexpr size_t bufsize = numeric_limits<V>::digits10 + 2;
    char buf[bufsize];
    const auto res = to_chars(buf, buf + bufsize, v);
    return S(buf, res.ptr);
}
```

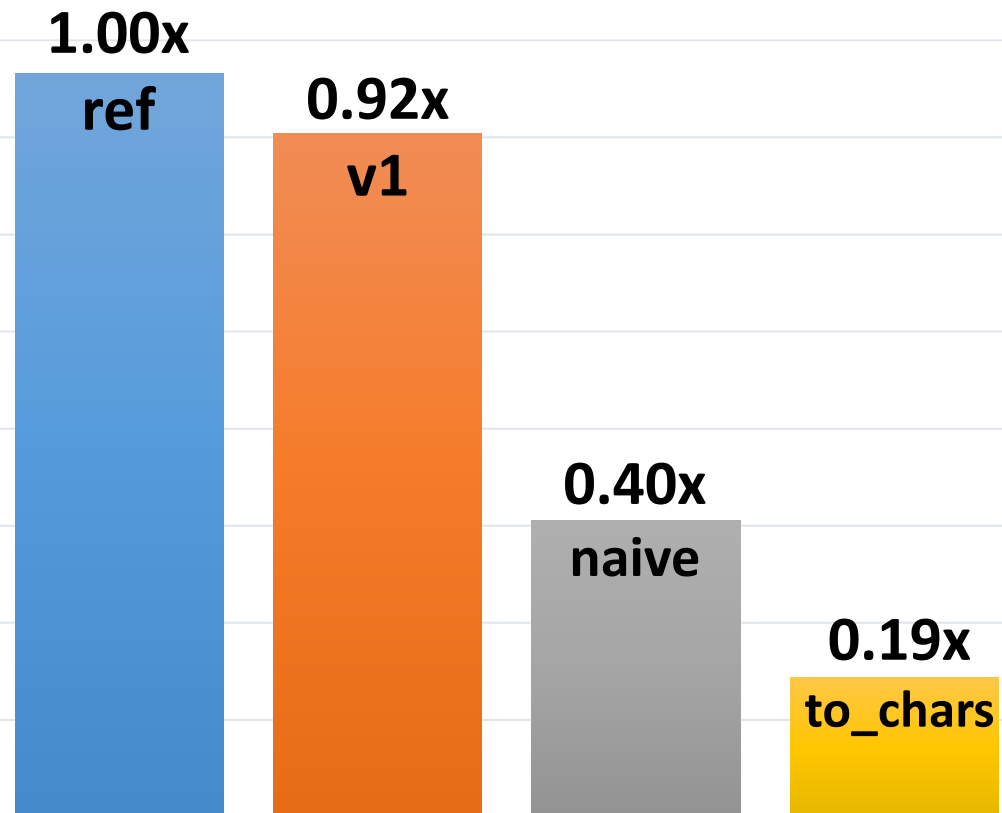
The value of `std::numeric_limits<T>::digits10` is the number of base-10 digits that can be represented by the type `T` without change.

`std::numeric_limits<std::uint8_t>::digits10 == 2:`  
any number in `[0, 99]` can be represented as `std::uint8_t` and `[256, 999]` can not.

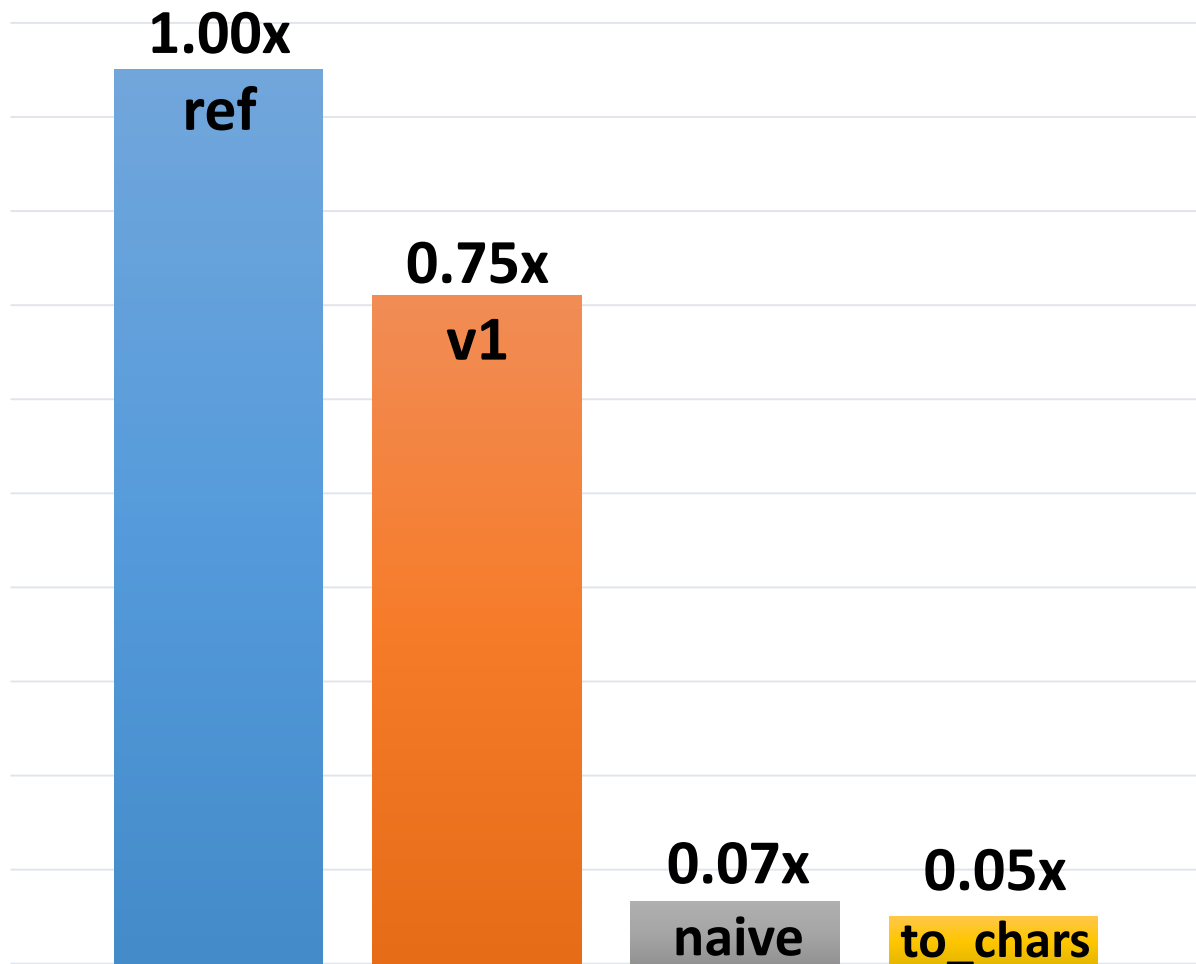
*to\_string(1)*



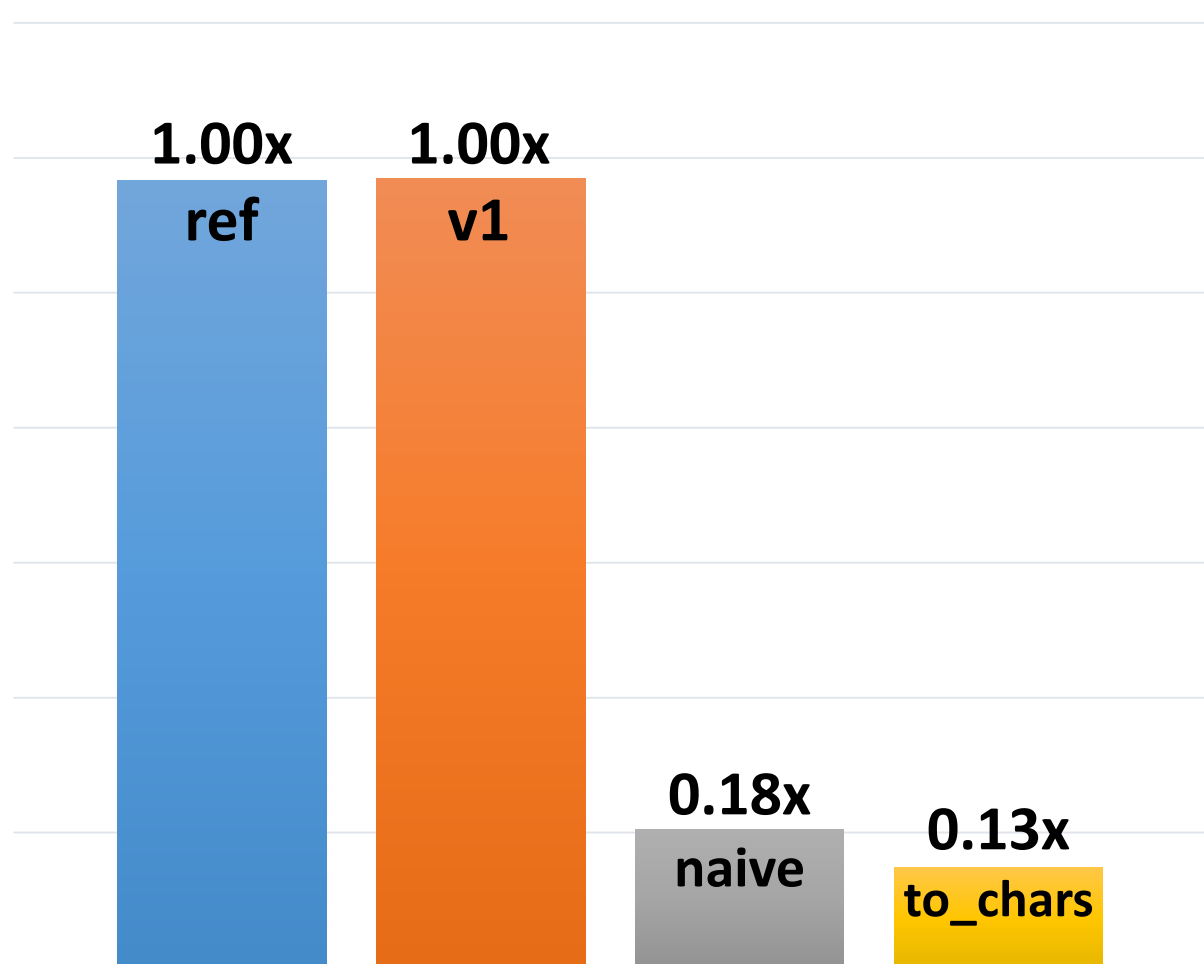
*to\_string(max)*



*to\_wstring(1)*



*to\_wstring(max)*



*std::to\_chars*



# divisions count

- naïve algorithm executes 2 divisions per character

# divisions count

- naïve algorithm executes 2 divisions per character
- who cares?

# divisions count

- naïve algorithm executes 2 divisions per character
- who cares? *idiv* instruction cost:

arch	cost (cycles) 32 bit	cost (cycles) 64 bit

<https://gmplib.org/~tege/x86-timing.pdf>

# divisions count

- naïve algorithm executes 2 divisions per character
- who cares? *idiv* instruction cost:

arch	cost (cycles) 32 bit	cost (cycles) 64 bit
Intel Core 2	40	116

<https://gmplib.org/~tege/x86-timing.pdf>

# divisions count

- naïve algorithm executes 2 divisions per character
- who cares? *idiv* instruction cost:

arch	cost (cycles) 32 bit	cost (cycles) 64 bit
Intel Core 2	40	116
Intel Nehalem	26	89

<https://gmplib.org/~tege/x86-timing.pdf>

# divisions count

- naïve algorithm executes 2 divisions per character
- who cares? *idiv* instruction cost:

arch	cost (cycles) 32 bit	cost (cycles) 64 bit
Intel Core 2	40	116
Intel Nehalem	26	89
Intel Sandy Bridge	26	92

<https://gmlib.org/~tege/x86-timing.pdf>

# divisions count

- naïve algorithm executes 2 divisions per character
- who cares? *idiv* instruction cost:

arch	cost (cycles) 32 bit	cost (cycles) 64 bit
Intel Core 2	40	116
Intel Nehalem	26	89
Intel Sandy Bridge	26	92
AMD K10	45	77

<https://gmplib.org/~tege/x86-timing.pdf>

# divisions count

- naïve algorithm executes 2 divisions per character
- who cares? *idiv* instruction cost:

arch	cost (cycles) 32 bit	cost (cycles) 64 bit
Intel Core 2	40	116
Intel Nehalem	26	89
Intel Sandy Bridge	26	92
AMD K10	45	77
Intel Atom	50	191

<https://gmlib.org/~tege/x86-timing.pdf>



# divisions count

- naïve algorithm executes 2 divisions per character
- who cares? *idiv* instruction cost:

*uint64\_t numbers processing  
should fall into 32-bit  
arithmetic!*

cost (cycles) 32 bit	cost (cycles) 64 bit
40	116
26	89
26	92
45	77
50	191

<https://gmplib.org/~tege/x86-timing.pdf>

# divisions count

- naïve algorithm executes 2 divisions per character
- who cares? *idiv* instruction cost:

*uint64\_t numbers processing  
should fall into 32-bit  
arithmetic!*

*(actually, no, wait a bit)*

cost (cycles) 32 bit	cost (cycles) 64 bit
40	116
26	89
26	92
45	77
50	191

<https://gmpilib.org/~tege/x86-timing.pdf>

# divisions count

- naïve algorithm executes 2 divisions per character
- *std::to\_chars* executes 1 division per character

```
static char digitLuts[201] =  
    "00010203040506070809"  
    "10111213141516171819"  
    "20212223242526272829"  
    "30313233343536373839"  
    "40414243444546474849"  
    "50515253545556575859"  
    "60616263646566676869"  
    "70717273747576777879"  
    "80818283848586878889"  
    "90919293949596979899";  
  
    i = val % 100;  
  
    digitLuts[2 * i]  
    digitLuts[2 * i + 1]  
  
    val /= 100
```

# divisions count

- naïve algorithm executes 2 divisions per character
- *std::to\_chars* executes 1 division per character

?

(compiler)

# divisions count

- naïve algorithm executes 2 divisions per character
- *std::to\_chars* executes 1 division per character

```
using T = uint64_t;
```

```
T f(T x, T y)
{
    return x / y;
}
```

```
pair<T, T> g(T x, T y)
{
    return { x / y, x % y };
}
```

# divisions count

- naïve algorithm executes 2 divisions per character
- *std::to\_chars* executes 1 division per character

```
using T = uint64_t;
```

```
T f(T x, T y)
{
    return x / y;
}
```

```
pair<T, T> g(T x, T y)
{
    return { x / y, x % y };
}
```

```
f(unsigned long, unsigned long)
    mov     rax, rdi
    xor     edx, edx
    div     rsi
    ret
```

```
g(unsigned long, unsigned long)
    mov     rax, rdi
    xor     edx, edx
    div     rsi
    ret
```

clang-9.0, -O2, x86

# divisions count

- naïve algorithm executes ~~2~~<sup>1</sup> divisions per character
- `std::to_chars` executes ~~1~~<sup>0,5</sup> division per character

# divisions count

- naïve algorithm executes 1 division per character
- *std::to\_chars* executes 0,5 divisions per character



# integers division

```
int f(int x)
{
    return x / 10;
}
```

```
int g(int x, int y)
{
    return x / y;
}
```

# integers division

```
int f(int x)
{
    return x / 10;
}
```

```
int g(int x, int y)
{
    return x / y;
}
```

```
f(int) : # @f(int)
    movsxd    rax, edi
    imul      rax, rax, 1717986919
    mov       rcx, rax
    shr       rcx, 63
    sar       rax, 34
    add       eax, ecx
    ret
```

```
g(int, int) : # @g(int, int)
    mov       eax, edi
    cdq
    idiv      esi
    ret
```

# integers division

```
int f(int x)
{
    return x / 10;
}
```

```
int g(int x, int y)
{
    return x / y;
}
```

```
f(int) : # @f(int)
    movsxd    rax, edi
    imul      rax, rax, 1717986919
    mov       rcx, rax
    shr       rcx, 63
    sar       rax, 34
    add       eax, ecx
    ret
```

```
g(int, int) : # @g(int, int)
    mov       eax, edi
    cdq
    idiv      esi
    ret
```

# integers division

```
int f(int x)
{
    return x / 10;
}
```

```
int g(int x, int y)
{
    return x / y;
}
```

```
f(int) : # @f(int)
        movsxd   rax, edi
        imul     rax, rax, 1717986919
        mov      rcx, rax
        shr      rcx, 63
        sar      rax, 34
        add      eax, ecx
        ret
```

```
g(int, int) : # @g(int, int)
        mov      eax, edi
        cdq
        idiv     esi
        ret
```

<https://libdivide.com/>

# integers division

```
int f(int x)
{
    return x / 10;
}
```

```
int g(int x, int y)
{
    return x / y;
}
```

```
f(int) : # @f(int)
        movsxd    rax, edi
        imul      rax, rax, 1717986919
        mov       rcx, rax
        shr       rcx, 63
        sar       rax, 34
        add       eax, ecx
        ret
```

```
g(int, int) : # @g(int, int)
        mov       eax, edi
        cdq
        idiv      esi
        ret
```

<https://lemire.me/blog/2019/02/08/faster-remainders-when-the-divisor-is-a-constant-beating-compilers-and-libdivide/>

<https://libdivide.com/>

# integers division

```
using T = std::uint64_t;
```

```
T f(T x)
{
    return x / 10;
}
```

```
pair<T, T> g(T x)
{
    return { x / 10, x % 10 };
}
```

# integers division

```
using T = std::uint64_t;
```

```
T f(T x)
{
    return x / 10;
}
```

```
pair<T, T> g(T x)
{
    return { x / 10, x % 10 };
}
```

```
f(unsigned long) :
mov     rax, rdi
movabs  rcx, -3689348814741910323
mul     rcx
mov     rax, rdx
shr     rax, 3
ret
```

```
g(unsigned long) :
movabs  rcx, -3689348814741910323
mov     rax, rdi
mul     rcx
mov     rax, rdx
shr     rax, 3
lea     rcx, [rax + rax]
lea     rcx, [rcx + 4 * rcx]
sub     rdi, rcx
mov     rdx, rdi
ret
```

# integers division

```
using T = std::uint64_t;
```

```
T f(T x)
{
    return x / 10;
}
```

```
f(unsigned long) :
    mov     rax, rdi
    movabs  rcx, -3689348814741910323
    mul     rcx
    mov     rax, rdx
    shr     rax, 3
    ret
```

```
pair<T, T> g(T x)
{
    return { x / 10, x % 10 };
}
```

```
g(unsigned long) :
    movabs  rcx, -3689348814741910323
    mov     rax, rdi
    mul     rcx
    mov     rax, rdx
    shr     rax, 3
    lea     rcx, [rax + rax]
    lea     rcx, [rcx + 4 * rcx]
    sub     rdi, rcx
    mov     rdx, rdi
    ret
```



# divisions count

- naïve algorithm executes 1 division per character
- *std::to\_chars* executes 0,5 divisions per character

# divisions count

optimized integer “divmod” op

- naïve algorithm executes 1 ~~division~~ per character
- `std::to_chars` executes 0,5 ~~divisions~~ per character

optimized integer “divmod” op

# divisions count

- naïve algorithm executes 1 optimized divmod per character
- *std::to\_chars* executes 0,5 optimized divmod per character

# length detection

```
if      (val < 10)
  ...
else if (val < 100)
  ...
else if (val < 1000)
  ...
else if (val < 10000)
  ...
else if (val < 100000)
  ...
else if (val < 1000000)
  ...
else if (val < 10000000)
```

# memcpy

```
static char digitLuts[201] =  
    "00010203040506070809"  
    "10111213141516171819"  
    "20212223242526272829"  
    "30313233343536373839"  
    "40414243444546474849"  
    "50515253545556575859"  
    "60616263646566676869"  
    "70717273747576777879"  
    "80818283848586878889"  
    "90919293949596979899";
```

```
char* append2(char* buffer, std::uint32_t i)  
{  
    std::memcpy(buffer, &digitLuts[(i) * 2], 2);  
    return buffer + 2;  
}
```

# memcpy

```
static char digitLuts[201] =  
    "00010203040506070809"  
    "10111213141516171819"  
    "20212223242526272829"  
    "30313233343536373839"  
    "40414243444546474849"  
    "50515253545556575859"  
    "60616263646566676869"  
    "70717273747576777879"  
    "80818283848586878889"  
    "90919293949596979899";
```

```
char* append2(char* buffer, std::uint32_t i)  
{  
    std::memcpy(buffer, &digitLuts[(i) * 2], 2);  
    return buffer + 2;  
}
```

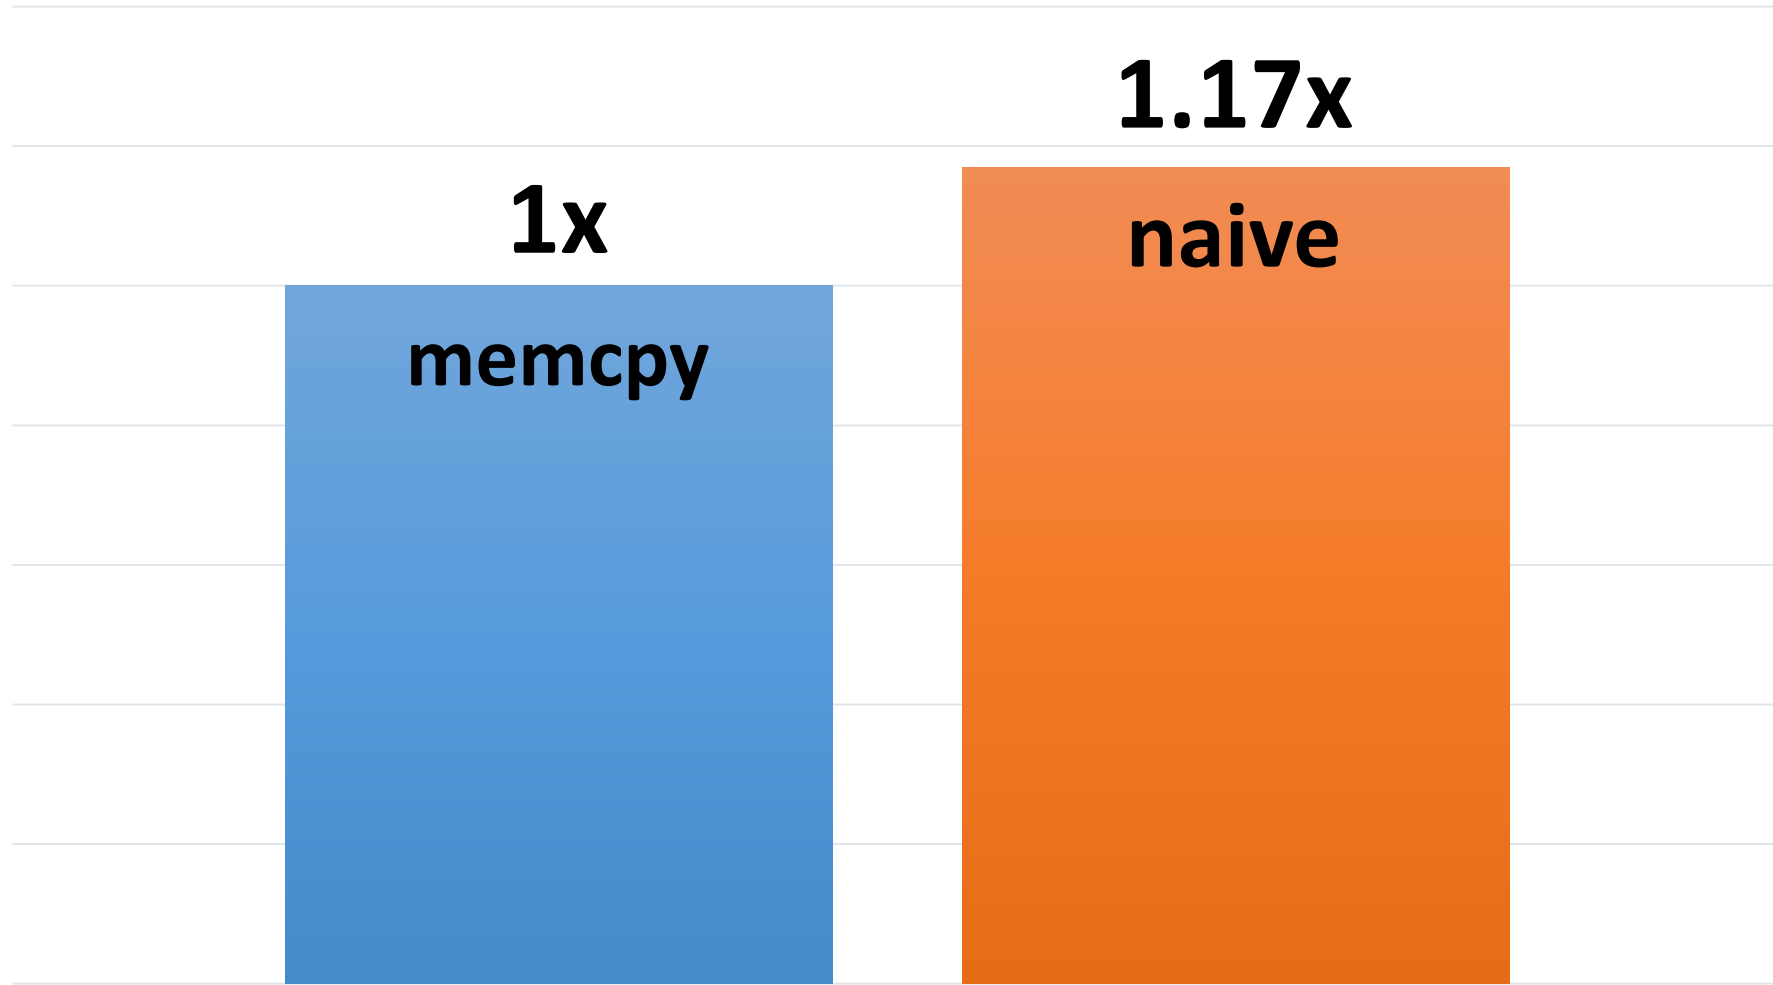
# memcpy

```
static char digitLuts[201] =  
    "00010203040506070809"  
    "10111213141516171819"  
    "20212223242526272829"  
    "30313233343536373839"  
    "40414243444546474849"  
    "50515253545556575859"  
    "60616263646566676869"  
    "70717273747576777879"  
    "80818283848586878889"  
    "90919293949596979899";
```

```
char* append2(char* buffer, std::uint32_t i)  
{  
    std::memcpy(buffer, &digitLuts[(i) * 2], 2);  
    return buffer + 2;  
}
```

```
char* append2(char* buffer, std::uint32_t i)  
{  
    *buffer = digitLuts[i * 2];  
    *(buffer + 1) = digitLuts[i * 2 + 1];  
    return buffer + 2;  
}
```

*to\_string(max)*





# memcpy

```
namespace ref
{
    char* append2(char* buffer, std::uint32_t i)
    {
        std::memcpy(buffer, &cDigitsLut[(i) * 2], 2);
        return buffer + 2;
    }
}
```

```
namespace tgt
{
    char* append2(char* buffer, std::uint32_t i)
    {
        *buffer = cDigitsLut[i * 2];
        *(buffer + 1) = cDigitsLut[i * 2 + 1];
        return buffer + 2;
    }
}
```

# memcpy

```
namespace ref
{
    char* append2(char* buffer, std::uint32_t i)
    {
        std::memcpy(buffer, &cDigitsLut[(i) * 2], 2);
        return buffer + 2;
    }
}
```

```
namespace tgt
{
    char* append2(char* buffer, std::uint32_t i)
    {
        *buffer = cDigitsLut[i * 2];
        *(buffer + 1) = cDigitsLut[i * 2 + 1];
        return buffer + 2;
    }
}
```

```
ref::append2(char*, unsigned int)
    add     esi, esi
    movzx   eax, word ptr[rsi + cDigitsLut]
    mov     word ptr[rdi], ax
    lea     rax, [rdi + 2]
    ret
```

```
tgt::append2(char*, unsigned int)
    add     esi, esi
    mov     al, byte ptr[rsi + cDigitsLut]
    mov     byte ptr[rdi], al
    mov     al, byte ptr[rsi + cDigitsLut + 1]
    mov     byte ptr[rdi + 1], al
    lea     rax, [rdi + 2]
    ret
```

# memcpy

```
namespace ref
{
    char* append2(char* buffer, std::uint32_t i)
    {
        std::memcpy(buffer, &cDigitsLut[(i) * 2], 2);
        return buffer + 2;
    }
}
```

```
namespace tgt
{
    char* append2(char* buffer, std::uint32_t i)
    {
        *buffer = cDigitsLut[i * 2];
        *(buffer + 1) = cDigitsLut[i * 2 + 1];
        return buffer + 2;
    }
}
```

```
ref::append2(char*, unsigned int)
    add     esi, esi
    movzx   eax, word ptr[rsi + cDigitsLut]
    mov     word ptr[rdi], ax
    lea     rax, [rdi + 2]
    ret
```

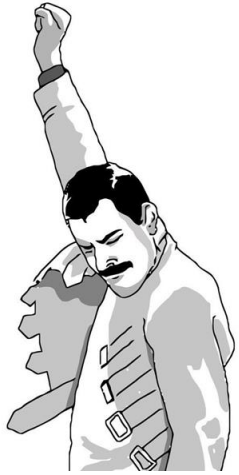
```
tgt::append2(char*, unsigned int)
    add     esi, esi
    mov     al, byte ptr[rsi + cDigitsLut]
    mov     byte ptr[rdi], al
    mov     al, byte ptr[rsi + cDigitsLut + 1]
    mov     byte ptr[rdi + 1], al
    lea     rax, [rdi + 2]
    ret
```

95% speedup... seems nice?

2 months later...

MERGED!

# MERGED!





**vlad.tsyrklevich** added a subscriber: **vlad.tsyrklevich**.

Jun 6 2019, 12:40 AM



After this change landed I started getting odd failures with check-llvm with MSan or ASan like the following:

[http://lab.llvm.org:8011/builders/sanitizer-x86\\_64-linux-bootstrap-msan/builds/12853](http://lab.llvm.org:8011/builders/sanitizer-x86_64-linux-bootstrap-msan/builds/12853)

/b/sanitizer-x86\_64-linux-bootstrap-msan/build/llvm/test/ThinLTO/X86/dot-dumper-full-lto.ll:12:10: error: CHECK: expected string not found in input

; CHECK: subgraph cluster\_4294967295

<stdin>:3:2: note: possible intended match here

subgraph cluster\_0004294967295 {



... and reverted in 5 hrs





 **vlad.tsyrklevich** added a subscriber: **vlad.tsyrklevich**.

Jun 6 2019, 12:40 AM

After this change landed I started getting odd failures with check-llvm with MSan or ASan like the following:

[http://lab.llvm.org:8011/builders/sanitizer-x86\\_64-linux-bootstrap-msan/builds/12853](http://lab.llvm.org:8011/builders/sanitizer-x86_64-linux-bootstrap-msan/builds/12853)

/b/sanitizer-x86\_64-linux-bootstrap-msan/build/llvm/test/ThinLTO/X86/dot-dumper-full-lto.ll:12:10: error: CHECK: expected string not found in input

; CHECK: subgraph cluster\_4294967295

<stdin>:3:2: note: possible intended match here

subgraph cluster\_0004294967295 {



... and reverted in 5 hrs:

clang checks failed in asan/msan mode

 **vlad.tsyrklevich** added a subscriber: **vlad.tsyrklevich**.

Jun 6 2019, 12:40 AM

After this change landed I started getting odd failures with check-llvm with MSan or ASan like the following:

[http://lab.llvm.org:8011/builders/sanitizer-x86\\_64-linux-bootstrap-msan/builds/12853](http://lab.llvm.org:8011/builders/sanitizer-x86_64-linux-bootstrap-msan/builds/12853)

/b/sanitizer-x86\_64-linux-bootstrap-msan/build/llvm/test/ThinLTO/X86/dot-dumper-full-lto.ll:12:10: error: CHECK: expected string not found in input

; CHECK: subgraph cluster\_4294967295

<stdin>:3:2: note: possible intended match here

subgraph cluster\_0004294967295 {



... and reverted in 5 hrs:

clang checks failed in asan/msan mode

*to\_string((uint64\_t)0xffffffff) == "0004294967295"*

# leading zeros problem

`std::to_chars (uint64_t)` adds redundant leading zeros for specific range of values

# leading zeros problem

`std::to_chars (uint64_t)` adds redundant leading zeros for specific range of values

Converts `value` into a character string by successively filling the range `[first, last)`, where `[first, last)` is required to be a valid range.

- 1) Integer formatters: `value` is converted to a string of digits in the given base (with no redundant leading zeroes). Digits in the range 10..35 (inclusive) are represented as lowercase characters a..z. If `value` is less than zero, the representation starts with a minus sign. The library provides overloads for all signed and unsigned integer types and for the type `char` as the type of the parameter `value`.

# leading zeros problem

`std::to_chars (uint64_t)` adds redundant leading zeros for specific range of values

Converts `value` into a character string by successively filling the range `[first, last)`, where `[first, last)` is required to be a valid range.

- 1) Integer formatters: `value` is converted to a string of digits in the given base (with no redundant leading zeroes). Digits in the range `10..35` (inclusive) are represented as lowercase characters `a..z`. If `value` is less than zero, the representation starts with a minus sign. The library provides overloads for all signed and unsigned integer types and for the type `char` as the type of the parameter `value`.

`to_chars` can puts leading zeros on numbers in  $[10^9, 10^{12})$

[https://bugs.llvm.org/show\\_bug.cgi?id=42166](https://bugs.llvm.org/show_bug.cgi?id=42166)

# leading zeros problem

`std::to_chars (uint64_t)` adds redundant leading zeros for specific range of values

Converts `value` into a character string by successively filling the range `[first, last)`, where `[first, last)` is required to be a valid range.

- 1) Integer formatters: `value` is converted to a string of digits in the given base (with no redundant leading zeroes). Digits in the range 10..35 (inclusive) are represented as lowercase characters a..z. If `value` is less than zero, the representation starts with a minus sign. The library provides overloads for all signed and unsigned integer types and for the type `char` as the type of the parameter `value`.

`to_chars` can puts leading zeros on numbers in  $[10^9, 10^{12})$

[https://bugs.llvm.org/show\\_bug.cgi?id=42166](https://bugs.llvm.org/show_bug.cgi?id=42166)

fix leading zeros in `std::to_chars`

<https://reviews.llvm.org/D63047>

MERGED!





# floating point numbers

<code>std::string to_string( float value );</code>	(7)	(since C++11)
<code>std::string to_string( double value );</code>	(8)	(since C++11)
<code>std::string to_string( long double value );</code>	(9)	(since C++11)

# floating point numbers

<code>std::string to_string( float value );</code>	(7)	(since C++11)
----------------------------------------------------	-----	---------------

<code>std::string to_string( double value );</code>	(8)	(since C++11)
-----------------------------------------------------	-----	---------------

<code>std::string to_string( long double value );</code>	(9)	(since C++11)
----------------------------------------------------------	-----	---------------

7,8) Converts a floating point value to a string with the same content as what `std::sprintf(buf, "%f", value)` would produce for sufficiently large buf.

9) Converts a floating point value to a string with the same content as what `std::sprintf(buf, "%Lf", value)` would produce for sufficiently large buf.

# floating point numbers

<code>std::string to_string( float value );</code>	(7)	(since C++11)
----------------------------------------------------	-----	---------------

<code>std::string to_string( double value );</code>	(8)	(since C++11)
-----------------------------------------------------	-----	---------------

<code>std::string to_string( long double value );</code>	(9)	(since C++11)
----------------------------------------------------------	-----	---------------

7,8) Converts a floating point value to a string with the same content as what `std::sprintf(buf, "%f", value)` would produce for sufficiently large buf.

9) Converts a floating point value to a string with the same content as what `std::sprintf(buf, "%Lf", value)` would produce for sufficiently large buf.

## Problems:

- *sprintf* depends on locale, *to\_chars* is locale independent!
- *to\_chars* for floating point numbers is not implemented yet
- *to\_chars* guarantees precise value recovery, *to\_string* does not

# floating point numbers

---

<code>std::string to_string( float value );</code>	(7)	(since C++11)
----------------------------------------------------	-----	---------------

---

<code>std::string to_string( double value );</code>	(8)	(since C++11)
-----------------------------------------------------	-----	---------------

---

<code>std::string to_string( long double value );</code>	(9)	(since C++11)
----------------------------------------------------------	-----	---------------

---

7,8) Converts a floating point value to a string with the same content as what `std::sprintf(buf, "%f", value)` would produce for sufficiently large buf.

9) Converts a floating point value to a string with the same content as what `std::sprintf(buf, "%Lf", value)` would produce for sufficiently large buf.

## Problems:

- *sprintf* depends on locale, *to\_chars* is locale independent!
- *to\_chars* for floating point numbers is not implemented yet
- *to\_chars* guarantees precise value recovery, *to\_string* does not

use proposal 1 to speedup *std::to\_string* for floating point numbers (1.0x – 5.2x):

<https://reviews.llvm.org/D64341>

# floating point numbers

---

```
std::string to_string( float value );
```

 (7) (since C++11)

---

```
std::string to_string( double value );
```

 (8) (since C++11)

---

```
std::string to_string( long double value );
```

 (9) (since C++11)

7,8) Converts a floating point value to a string with the same content as what `std::sprintf(buf, "%f", value)` would produce for sufficiently large buf.

9) Converts a floating point value to a string with the same content as what `std::sprintf(buf, "%Lf", value)` would produce for sufficiently large buf.

## Problems:

- *sprintf* depends on locale, *to\_chars* is locale independent!
- *to\_chars* for floating point numbers is not implemented yet
- *to\_chars* guarantees precise value recovery, *to\_string* does not

use proposal 1 to speedup *std::to\_string* for floating point numbers (1.0x – 5.2x):

<https://reviews.llvm.org/D64341>

# libstdc++(GNU) and MS STL

	then	now (23 nov. 2019)
libstdc++(GNU)		
MS STL		

# libstdc++(GNU) and MS STL

	then	now (23 nov. 2019)
libstdc++(GNU)	proposal 1 success path ( <i>sprintf</i> )	
MS STL	-	

# libstdc++(GNU) and MS STL

	then	now (23 nov. 2019)
libstdc++(GNU)	proposal 1 success path ( <i>sprintf</i> )	<i>to_chars</i>
MS STL	-	naïve algorithm



# Results



*to\_string* / *to\_wstring* performance  
improved up to 20x times

<https://reviews.llvm.org/D59178>



fixed leading zeros in `std::to_chars`

<https://reviews.llvm.org/D63047>

Thank you